

# Permutation Tests

Livio Finos

University of Padova

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Introduction . . . . .	2
1.2	Renewed Interest in Permutation Testing . . . . .	2
1.3	The <code>flip</code> Package . . . . .	2
1.4	The Age vs Reaction Time Dataset . . . . .	3
1.5	Measuring Dependence Between Two Variables . . . . .	4
<b>2</b>	<b>Permutation Approach to Hypothesis Testing</b>	<b>6</b>
2.1	Permutation Tests in a Nutshell . . . . .	6
2.2	Summary . . . . .	11
2.3	A More Formal Approach . . . . .	13
2.4	Comparison with Parametric Linear Models . . . . .	15
2.5	Permutationally Equivalent Tests . . . . .	17
<b>3</b>	<b>Special Cases</b>	<b>18</b>
3.1	Rank Correlation . . . . .	18
3.2	Two Independent Samples Problem . . . . .	20
3.3	Chi-square and Other Categorical Methods . . . . .	21
3.4	ANOVA (C-sample) . . . . .	23
3.5	Stratified Permutations (Discrete Nuisances) . . . . .	25
<b>4</b>	<b>Paired samples</b>	<b>25</b>
4.1	Advantages of this approach . . . . .	26
4.2	Repeated measures and mixed models . . . . .	28
<b>5</b>	<b>Multivariate Testing</b>	<b>28</b>
5.1	Seeds Data . . . . .	28
5.2	Marginal vs Joint Distribution . . . . .	29
5.3	Rejection Regions (and Overall Testing) . . . . .	31

<b>6 FWER Control via Permutation Tests</b>	<b>34</b>
6.1 Permutation Bonferroni . . . . .	34
6.2 Improved Bonferroni . . . . .	34
6.3 Multiple Testing via Permutations . . . . .	35
6.4 P-value Correlation Structure . . . . .	35
6.5 Improved Holm: Westfall & Young . . . . .	35
6.6 General Framework: Closed Testing . . . . .	36
<b>7 Case Study: Pharmacokinetic Study of Carbidopa</b>	<b>37</b>
7.1 Solution . . . . .	38
<b>8 Minimal Bibliography</b>	<b>42</b>

# 1 Introduction

## 1.1 Introduction

- Well-established nonparametric inference approach: Fisher (1935); Pitman (1937, 1938).
- Generally requires fewer assumptions about the data generating process than parametric counterparts.
- Excellent inferential properties, typically:
  - Exactness (exact control of Type I error)
  - Asymptotic optimality and convergence to parametric counterparts when they exist.
- Fisher’s exact test is a prototypical example, but
- The general approach had limited applicability without computational support.

## 1.2 Renewed Interest in Permutation Testing

- A milestone: Westfall and Young (1993). *Resampling-Based Multiple Testing: Examples and Methods for p-value Adjustment*. Wiley.
- Many active research areas adopt these methods in daily statistical analysis (e.g., genetics and neuroscience: Nichols and Holmes (2002); Pantazis et al. (2009); Winkler et al. (2014)).
- Permutation approach:
  - Ideal for **randomized experimental designs**
  - Handles complex models without formal definition of the data generating process.

## 1.3 The **flip** Package

Available on CRAN and GitHub (<https://github.com/livioivil/flip>).

To install the GitHub version in R:

```
library(devtools)
install_github('liviocivil/flip')
```

Before starting

```
# Clean memory
rm(list = ls())

# Customize graph output
par.old <- par()
par(cex.main = 1.5, lwd = 2, col = "darkgrey", pch = 20, cex = 3)
palette(c("#FF0000", "#00A08A", "#FFCC00", "#445577", "#45abff"))

# Customize knitr output
knitr::opts_chunk$set(fig.align = "center") # fig.width=6, fig.height=6
```

## 1.4 The Age vs Reaction Time Dataset

Subjects' reaction times were tested by having them grab a meter stick after it was released. The number of centimeters the meter stick dropped before being caught directly measures response time.

Age values are in years. Gender is coded as F for female and M for male. Reaction.Time values are in centimeters.

(Data are fictitious)

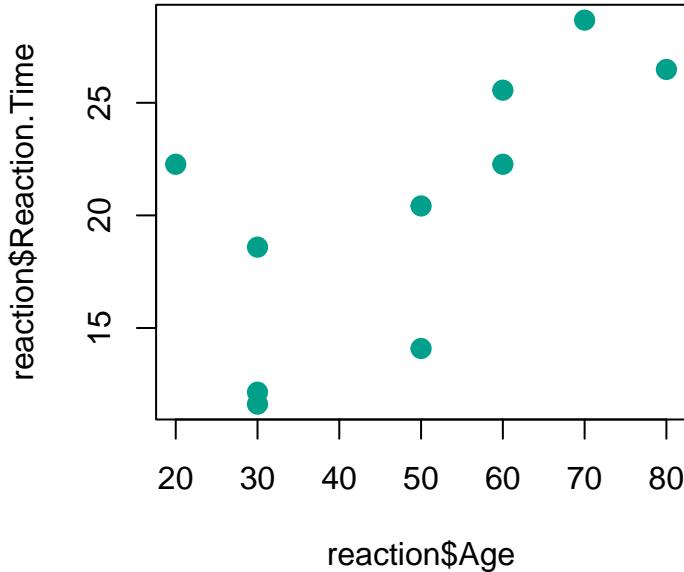
To read the data:

```
data(reaction, package = "flip")
# Alternatively, download from: https://github.com/liviocivil/flip/tree/master/data
# or load("reaction.rda")
str(reaction)

## Classes 'tbl_df', 'tbl' and 'data.frame':   10 obs. of  3 variables:
## $ Age          : num  70 50 30 60 80 60 30 30 20 50
## $ Gender       : Factor w/ 2 levels "F","M": 1 1 2 1 2 1 2 2 1 2
## $ Reaction.Time: num  28.7 20.4 11.6 22.3 26.5 ...
```

Plot the data:

```
plot(x = reaction$Age, y = reaction$Reaction.Time, pch = 20, col = 2, cex = 2)
```



## 1.5 Measuring Dependence Between Two Variables

Define:  
-  $X = \text{Age}$   
-  $Y = \text{Reaction.Time}$

Review common indices for measuring (linear) dependence between two variables.

### 1.5.1 Covariance and Variance

**Covariance** between  $X$  and  $Y$ :

$$\sigma_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{n}$$

- Values between  $-\infty$  and  $\infty$
- $\sigma_{xy} \approx 0$ : No dependency between  $X$  and  $Y$
- $\sigma_{xy} \gg 0$  ( $\ll 0$ ): Strong positive (negative) dependency

**Variance** of  $X$ :

$$\sigma_{xx} = \sigma_x^2 = \frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n}$$

**Standard Deviation** of  $X$ :

$$\sigma_x = \sqrt{\sigma_{xx}}$$

### 1.5.2 Correlation

Covariance alone makes it difficult to assess relationship strength. Note that:

$$-\sigma_x \sigma_y \leq \sigma_{xy} \leq \sigma_x \sigma_y$$

is equivalent to:

$$-1 \leq \frac{\sigma_{xy}}{\sigma_x \sigma_y} \leq 1$$

**Correlation** between  $X$  and  $Y$ :

$$\rho_{xy} = \frac{\sigma_{xy}}{\sigma_x \sigma_y} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

- Values between  $-1$  and  $1$
- $\rho_{xy} \approx 0$ : No dependency
- $\rho_{xy} \approx 1$  ( $-1$ ): Strong positive (negative) dependency

### 1.5.3 Linear Trend: Least Squares Method

Describe the relationship between `Reaction.Time` and `Age` with a straight line:

$$E(Reaction.Time) \approx \beta_0 + \beta_1 Age$$

$$E(Y) = \beta_0 + \beta_1 X$$

Draw a line through the center of the data.

**Least-squares estimator:** Find parameters minimizing the sum of squared residuals:

Find  $\hat{\beta}_0$  and  $\hat{\beta}_1$  that minimize:  $\sum_{i=1}^n (y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i))^2$

Estimates:

- Slope:  $\hat{\beta}_1 = \frac{\sigma_{xy}}{\sigma_{xx}} = \rho_{xy} \frac{\sigma_y}{\sigma_x} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = 0.2064719$
- Intercept:  $\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = 10.3013483$  - Estimated response:  $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$  - Residuals:  $y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i) = y_i - \hat{y}_i$

Sum of squared residuals:  $\sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2$

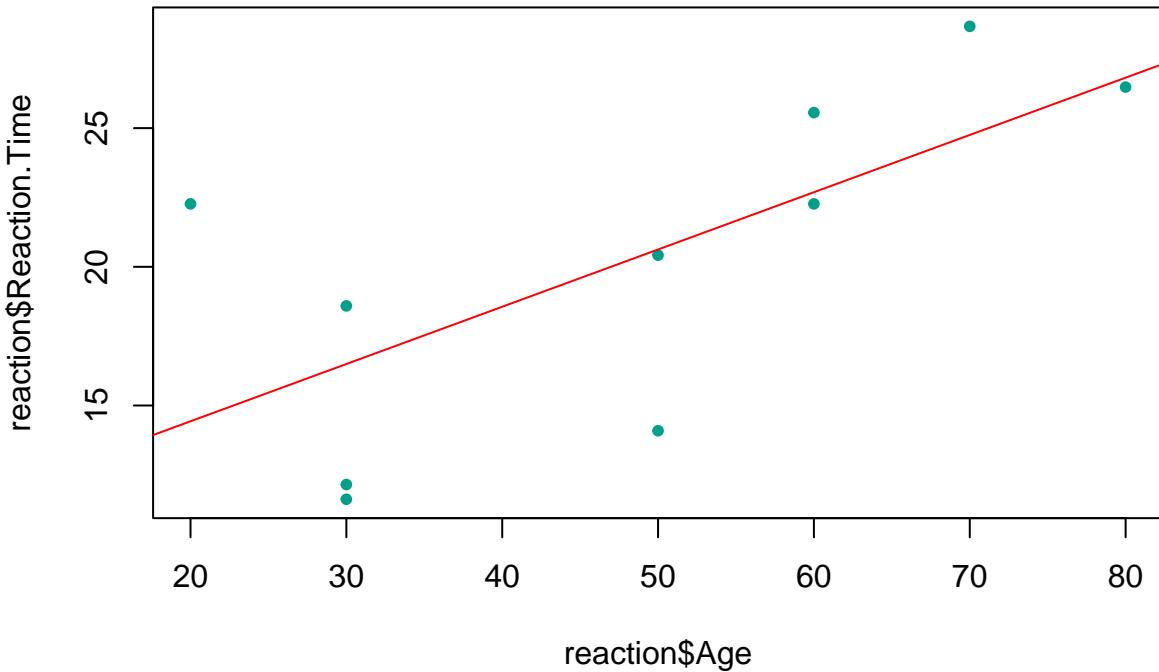
Visual representation:

```
model <- lm(Reaction.Time ~ Age, data = reaction)
coefficients(model)

## (Intercept)          Age
## 10.3013483  0.2064719

plot(reaction$Age, reaction$Reaction.Time, pch = 20, col = 2, cex = 1)
coeff <- round(coefficients(model), 1)
title(paste("Y =", coeff[1], "+", coeff[2], "* X"))
abline(model, col = 1)
```

$$Y = 10.3 + 0.2 * X$$



## 2 Permutation Approach to Hypothesis Testing

### 2.0.1 Preliminary Remarks

Note that all measures above make no assumptions about the random process generating the data.

Now assume  $Y$  (and possibly  $X$ ) is generated by a random variable. Further minimal assumptions will be specified later.

**Question: Is there a relationship between  $Y$  and  $X$ ?**

We estimated  $\hat{\beta}_1 = 0.2064719$

But is the **true value**  $\beta_1$  actually different from 0 (indicating no relationship)? Or is the difference from 0 due to random sampling?

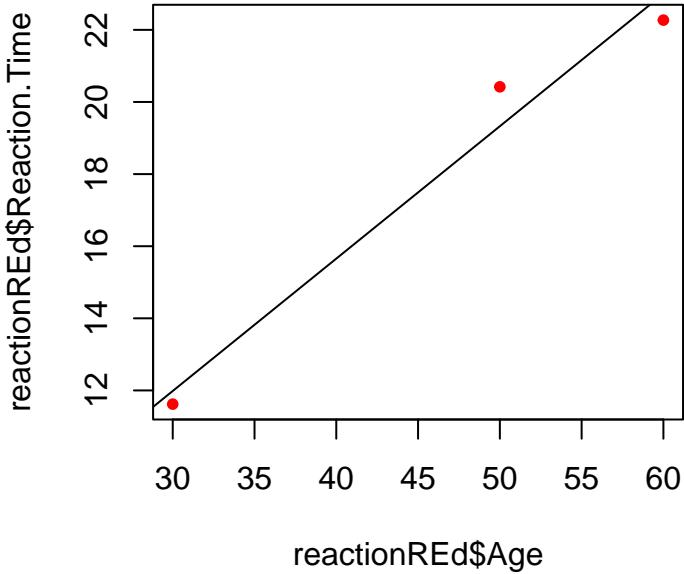
- **Null Hypothesis**  $H_0 : \beta_1 = 0$  (the **true**  $\beta_1$ , not its estimate  $\hat{\beta}_1$ ). No relationship between  $X$  and  $Y$ .
- **Alternative Hypothesis**  $H_1 : \beta_1 > 0$  (positive relationship).

Other possible  $H_1$  specifications:  $\beta_1 < 0$  or, more commonly,  $\beta_1 \neq 0$ .

### 2.1 Permutation Tests in a Nutshell

As a toy example, use a data subset:

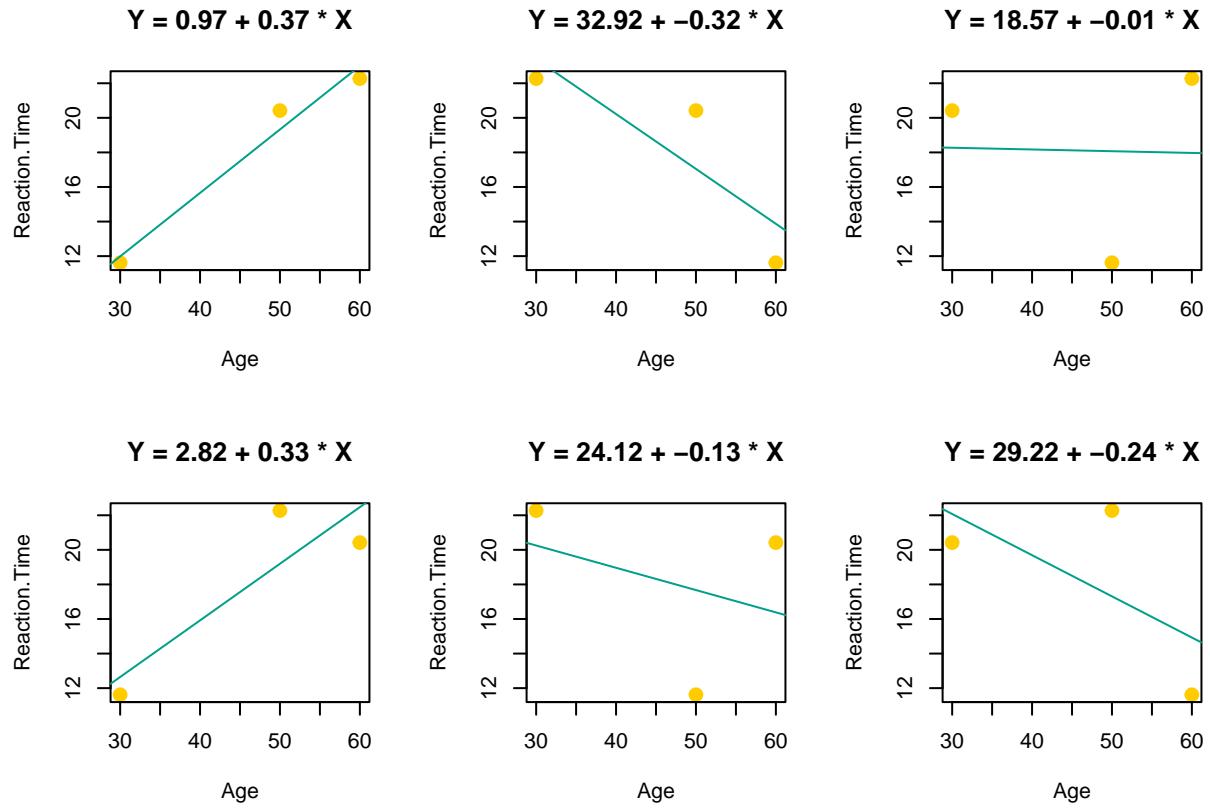
```
##   Age Gender Reaction.Time
## 2  50      F        20.42
## 3  30      M        11.62
## 4  60      F        22.27
```



- If  $H_0$  is true: No linear relationship between  $X$  and  $Y$
- Observed trend is due to chance
- Any other pairing of  $x_i$  and  $y_i$  was equally likely
- Generate hypothetical datasets by permuting  $Y$  observations
- How many equally likely datasets exist with observed  $X$  and  $Y$ ?  $3 \times 2 \times 1 = 3! = 6$  possible datasets.

**Remark:** We only assume  $Y$  is a random variable. The key assumption is exchangeability: the joint density  $f(y_1, \dots, y_n)$  is invariant to permutations of  $y_1, \dots, y_n$ .

### 2.1.1 All Potential Datasets



#### 2.1.1.1 In Our Complete Dataset

Apply the same principle to the full dataset...

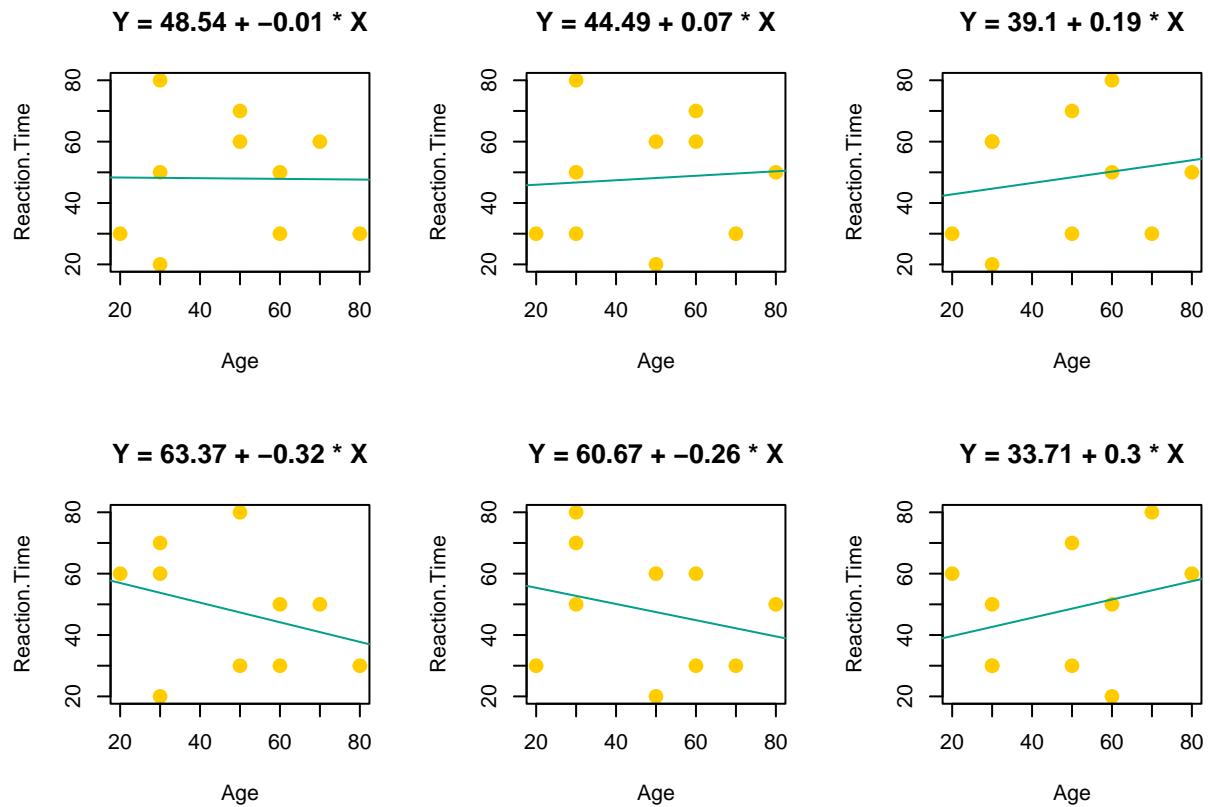
How many permutations of  $y_1, \dots, y_n$  are possible?  $n! = 10! = 3,628,800$ .

Manageable, but with  $n = 20$ ?  $20! = 2.43 \times 10^{18}$  - too large!

Calculate a smaller but sufficiently large number  $B$  of random permutations.

Examples:

**Age vs Permuted Reaction.Time**



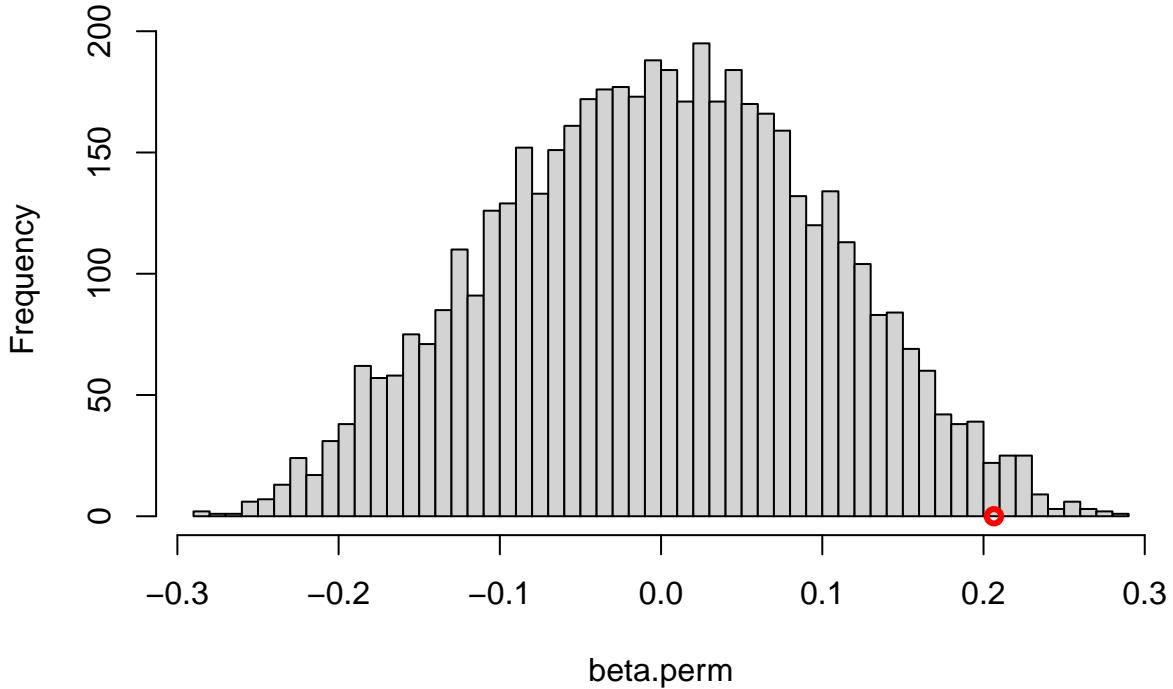
Repeat 5000 times and examine the  $\hat{\beta}_1$  histogram:

```
# beta_1 estimated on observed data
beta1 <- coefficients(lm(Reaction.Time ~ Age, data = reaction)) [2]

# Function to permute y and calculate beta_1
my.beta.perm <- function(Y, X) {
  model <- lm(sample(Y) ~ X)
  coefficients(model) [2]
}

# Replicate B-1 times
beta.perm <- replicate(B, my.beta.perm(reaction$Reaction.Time, reaction$Age))
```

**Histogram of beta.perm**



### 2.1.2 How Likely Was $\hat{\beta}_1^{obs}$ ?

(before the experiment!)

What was the probability of obtaining a value  $\geq \hat{\beta}_1^{obs}$  among possible  $\hat{\beta}_1^{*b}$  values (from permuted data)?

Remarks: -  $\hat{\beta}_1^{*b} < \hat{\beta}_1^{obs}$  (closer to 0): Less evidence against  $H_1$  than  $\hat{\beta}_1^{obs}$  -  $\hat{\beta}_1^{*b} \geq \hat{\beta}_1^{obs}$ : Equal or stronger evidence for  $H_1$  than  $\hat{\beta}_1^{obs}$

### 2.1.3 P-value Calculation

Out of  $B = 5000$  permutations, we obtained 4921 cases where  $\hat{\beta}_1^{*b} \leq \hat{\beta}_1^{obs}$ .

The p-value (significance) is:  $p = \frac{\#\{\hat{\beta}_1^{*b} \geq \hat{\beta}_1^{obs}\}}{B} = 0.0162$

( $\hat{\beta}_1^{obs}$  counts as one random permutation)

### 2.1.4 Interpretation

The probability  $P(\hat{\beta}_1^* \geq \hat{\beta}_1 = 0.206 | H_0) = p = 0.0162$  is very small. It was unlikely to obtain such a value if  $H_0$  is true.

The Neyman-Pearson approach established significance thresholds like  $\alpha = .05$  (or = .01). When  $p \leq \alpha$ , we reject  $H_0$  (no relationship). We then conclude  $H_1$  is likely true (positive relationship).

- Type I error: False Positive  
True hypothesis is  $H_0$  (no correlation), but we accept  $H_1$  (positive correlation)
- Type II error: False Negative  
True hypothesis is  $H_1$  (positive correlation), but we fail to reject  $H_0$  (no correlation)

## 2.2 Summary

**p-value:** Proportion of experiments providing equal or stronger evidence against  $H_0$  compared to observed data.

To compute it, we need: - **Orbit  $\mathcal{O}$**  - **Test statistic** ( $T : \mathbb{R}^n \rightarrow \mathbb{R}$ ) quantifying evidence against  $H_0$  - Higher values indicate stronger evidence against  $H_0$  - Computing  $T$  for each  $\mathcal{O}$  element induces an ordering on  $\mathcal{O}$

In our example:  $T = \hat{\beta}_1 = \hat{\sigma}_{xy}/\hat{\sigma}_{xx}$  (estimated slope). Higher slope indicates stronger evidence for  $H_1$ .

### Type I Error Control

We want to limit false discoveries (be conservative). Bound the probability of false discovery:

$$P(\text{p-value} \leq \alpha | H_0) \leq \alpha$$

This ensures that over many experiment replications, we find false correlations with probability  $\alpha$  (e.g.,  $0.05 = 5\%$ ).

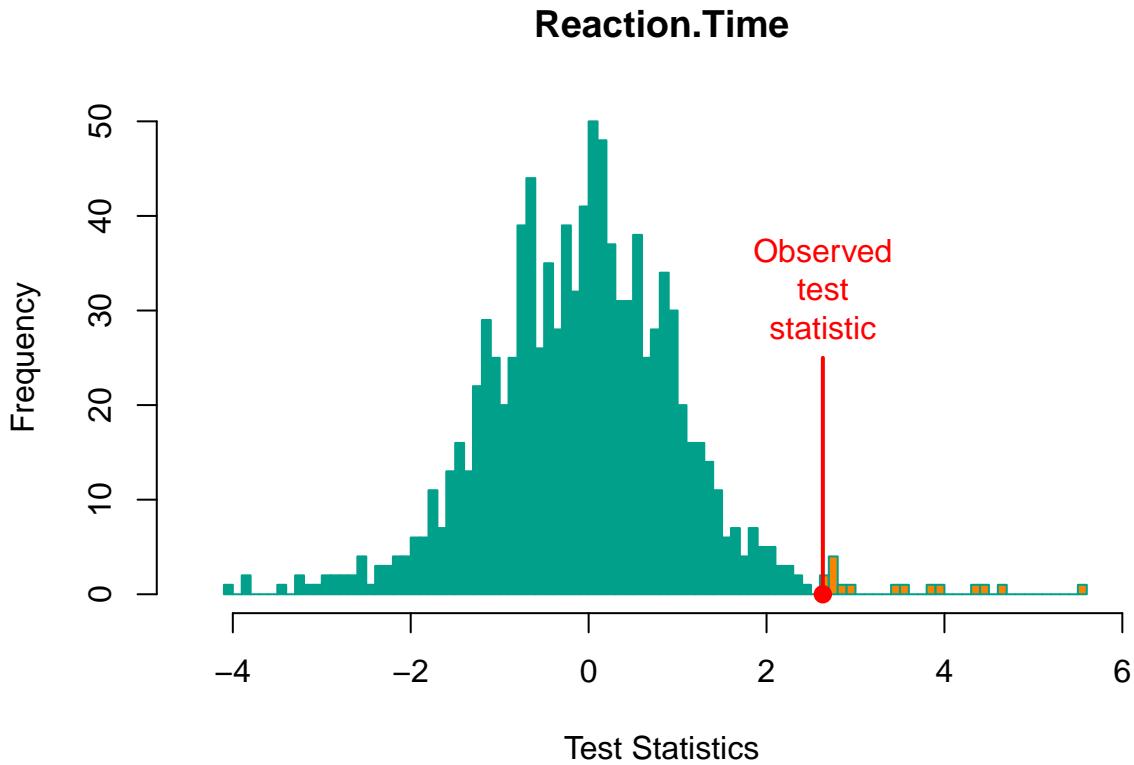
#### 2.2.1 Implementation in flip

```
library(flip)
(res <- flip(Reaction.Time ~ Age, data = reaction, tail = 1))

##
##           Test   Stat tail p-value
## Reaction.Time     t 2.633    > 0.0150

# Compare with:
# flip(Reaction.Time ~ Age, data = reaction, tail = 1, statTest = "cor")
# flip(Reaction.Time ~ Age, data = reaction, tail = 1, statTest = "coeff")

plot(res)
```



### Type I Error Control

We want to guarantee few false positives. Bound false discovery probability:

$$P(\text{p-value} \leq \alpha | H_0) \leq \alpha$$

This ensures long-run false correlation discovery rate  $\leq \alpha$  (e.g., 5%).

#### 2.2.2 Two-sided Alternatives

$H_1 : \beta_1 > 0$  (positive relationship) requires a priori justification.

More commonly, the two-sided alternative is appropriate:  $H_1 : \beta_1 \neq 0$  (relationship exists, direction unspecified)

We consider both very small and very large estimated coefficients as anomalous ('far from 0').

$$\text{P-value: } p = \frac{\#\{|\hat{\beta}_1^{*b}| \geq |\hat{\beta}_1^{obs}|\}}{B} = 0.0326$$

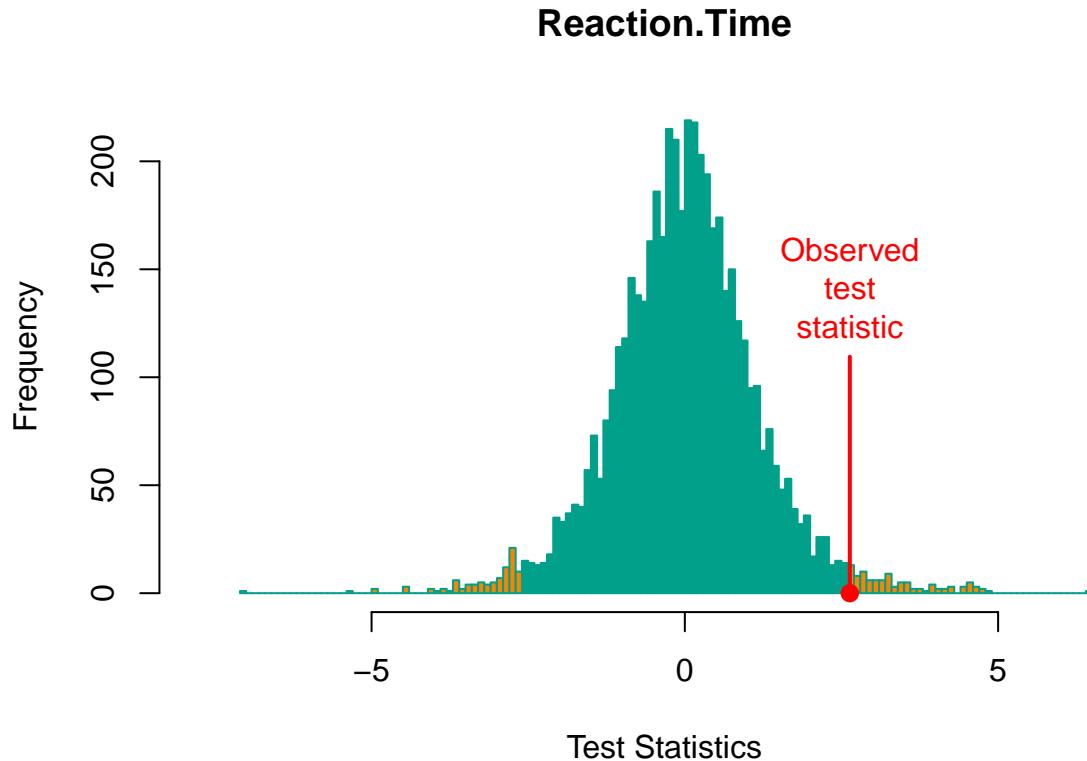
(Note: observed test statistic included among permutations)

In `flip`:

```
library(flip)
(res <- flip(Reaction.Time ~ Age, data = reaction, tail = 0, perms = 5000))
```

```
##                                     Test Stat tail p-value
## Reaction.Time      t 2.633    >< 0.0374
```

```
plot(res)
```



### 2.3 A More Formal Approach

(See also Pesarin, 2001; Hemerik & Goeman, 2018)

Let  $Y$  be data in sample space  $\mathcal{Y}$ . Let  $\Pi$  be a (usually finite) set of transformations  $\pi : \mathcal{Y} \rightarrow \mathcal{Y}$  forming a **group** under composition: - Contains identity - Every element has an inverse - Closed: if  $\pi_1, \pi_2 \in \Pi$ , then  $\pi_1 \circ \pi_2 \in \Pi$

(e.g.,  $\Pi =$  all possible permutations)

**Null Hypothesis**

$$H_0 : Y \sim P \in \Omega_0$$

**Randomization Hypothesis** Under  $H_0$ , the distribution of  $Y$  is invariant under  $\Pi$ : for every  $\pi \in \Pi$ ,  $\pi Y$  and  $Y$  have the same distribution when  $Y \sim P \in \Omega_0$ .

(See also Lehmann & Romano, 2006. *Testing Statistical Hypotheses*. Springer.)

**Test statistic**  $T(Y) : \mathbb{R}^n \rightarrow \mathbb{R}$

$T^{(k)}(Y)$  is the  $\lceil (1 - \alpha) |\Pi| \rceil$ -th sorted value of  $T(\pi Y)$

Define the test:

$$\phi(Y) = \begin{cases} 1 & \text{if } T(Y) > T^{(k)}(Y) \\ 0 & \text{otherwise} \end{cases}$$

**Theorem:** Under  $H_0$ ,  $E_P(\phi(Y)) = \alpha$ , i.e.,  $P(T(Y) > T^{(k)}) \leq \alpha$ .

**Proof**

For each  $Y \in \mathcal{Y}$ , define orbit  $O_Y = \{\pi Y : \pi \in \Pi\} \subseteq \mathcal{Y}$ .

Let  $A = \{Y \in \mathcal{Y} : T(Y) > T^{(k)}(Y)\}$  be the rejection set. Under  $H_0$ , by group structure,  $\Pi\pi = \Pi$  for all  $\pi \in \Pi$ , so  $T^{(k)}(\pi Y) = T^{(k)}(Y)$  for all  $\pi \in \Pi$ . Thus:

$$\#\{\pi \in \Pi : \pi Y \in A\} = \#\{\pi \in \Pi : T(\pi Y) > T^{(k)}(\pi Y)\} = \#\{\pi \in \Pi : T(\pi Y) > T^{(k)}(Y)\} \leq \alpha |O_Y|$$

Endow orbits with inherited  $\sigma$ -algebra. As in Lehmann (2005, Theorem 15.2.2):

$$P(Y \in A | O_Y) = \frac{|A|}{|O_Y|}$$

Bounded by  $\alpha$ . Hence:

$$P(Y \in A) = \mathbb{E}\{P(Y \in A | O_Y)\} \leq \alpha$$

### Alternative Proof

By construction,  $\sum_{\pi \in \Pi} \phi(\pi Y) = |\Pi|\alpha$ . Thus:

$$|\Pi|\alpha = E_P \left( \sum_{\pi \in \Pi} \phi(\pi Y) \right) = \sum_{\pi \in \Pi} E_P(\phi(\pi Y))$$

By null hypothesis:  $E_P(\phi(Y)) = E_P(\phi(\pi Y))$ , so:

$$|\Pi|\alpha = \sum_{\pi \in \Pi} E_P(\phi(Y)) = |\Pi|E_P(\phi(Y))$$

giving  $E_P(\phi(Y)) = \alpha$ .

### 2.3.1 Further Notes on Permutation Testing

**Orbit:**  $\mathcal{O} = \{\pi Y : \pi \in \Pi\} \subseteq \mathcal{Y}$

i.e.  $\mathcal{O} = \{\text{all permutations of observed data } \mathbf{y}\} = \{\mathbf{y}^* : \pi^* \circ \mathbf{y}\}$

**2.3.1.1 Exchangeability** **Exchangeability Assumption:** Under  $H_0$ , observations are exchangeable: e.g.,  $f(y_1, y_2) = f(y_2, y_1)$ .

Therefore the  $\mathcal{O}$  is a set of samples in  $\mathcal{Y}$  sharing the same likelihood under the null hypothesis:  $\mathcal{O} = \{\pi \mathbf{y} : f_{H_0}(\pi \mathbf{y}) = f_{H_0}(\mathbf{y})\}$

( $|\mathcal{O}| = \text{number of elements}$ )

With exchangeability:

**Proof Intuition** (alternative Type I error control proof):

$$f(\mathbf{y} | \mathcal{O}) = \frac{f(\mathbf{y} \cap \mathcal{O})}{f(\mathcal{O})} = \frac{f(\mathbf{y})}{f(\mathcal{O})} = \frac{f(\mathbf{y})}{f(\bigcup_{y \in \mathcal{O}} y)} = \frac{1}{|\mathcal{O}|} \quad \forall \mathbf{y} \in \mathcal{O}$$

Each permutation is equally likely in  $\mathcal{O}$  (due to group structure).

$$\begin{aligned} E(\phi(Y) | \mathbf{y} \in \mathcal{O}, H_0) &= P(T(\mathbf{y}) > T^{(k)} | \mathbf{y} \in \mathcal{O}, H_0) \\ &= \int_{T^{(k)}}^{+\infty} f(T(\mathbf{y})) dT(\mathbf{y}) \\ &= \sum_{\mathbf{y} \in \mathcal{O}} I(T(\mathbf{y}) > T^{(k)}) / |\mathcal{O}| \leq \alpha \quad \forall \mathcal{O} \end{aligned}$$

Then:  $E(\phi(\mathbf{y})) = \int_P E(\phi(\mathbf{y}) | \mathbf{y} \in \mathcal{O}, H_0) d\mathbf{y}$

**2.3.1.2 Independence vs exchangeability** Always true if observations:

- Are **identically distributed** - Have **same dependence** (e.g., same correlation)

Parametric  $t$ -tests and linear models assume independence (stricter than ‘same dependence’) and normality of errors—more stringent than permutation approach.

When normality fails, parametric approaches only provide asymptotic Type I error control, while permutation provides exact control.

### 2.3.2 Properties (see Pesarin, 2001)

The theorem proves permutation tests have **exact Type I error control**:  $P(\text{p-value} \leq \alpha | H_0) = \alpha$ , assuming  $\alpha \in \{1/|\mathcal{O}|, 2/|\mathcal{O}|, \dots, 1\}$  (since  $\mathcal{O}$  is finite and  $T(\pi y)$  distribution is a step function). For other  $\alpha$  values, tests are slightly conservative (or require randomized tests, not discussed here).

Additional properties: - **Unbiased**:  $P(\text{p-value} \leq \alpha | H_1) > \alpha$  - **Consistent**:  $P(\text{p-value} \leq \alpha | H_1) \rightarrow 1$  as  $n \rightarrow \infty$  - Converges to parametric counterpart when it exists

### 2.3.3 Estimated p-values

In practice, the p-value is often *estimated* using random permutations when it is computationally infeasible to compute the exact permutation p-value based on the entire permutation group.

Random permutations are typically drawn uniformly from the orbit  $\mathcal{O}$  without replacement. In this Section,  $p$  denotes the exact p-value computed from the entire orbit  $\mathcal{O}$ , and  $\hat{p}$  denotes its estimate computed from  $B$  randomly sampled permutations.

If we force the first element to be the observed test statistic  $T(Y)$ , then  $T(Y)$  becomes over-represented in the sample of  $B$  elements from  $\mathcal{O}$ . Consequently,  $E(\hat{p}) > p$ , although  $\lim_{B \rightarrow \infty} \hat{p} = p$ .

An unbiased estimator of  $p$ , denoted  $\hat{p}_0$ , can be obtained by removing the constraint that  $T(Y)$  must be included in the sample (i.e., by sampling permutations without including the observed statistic by default). This ensures  $E(\hat{p}_0) = p$ . However, as Phipson and Smyth (2010) thoroughly explain, using this unbiased estimator can be problematic because  $\hat{p}_0$  is almost never stochastically larger than a uniform distribution on  $[0, 1]$  under  $H_0$ . This is evident from the fact that  $\hat{p}_0$  typically has a positive probability of being exactly zero.

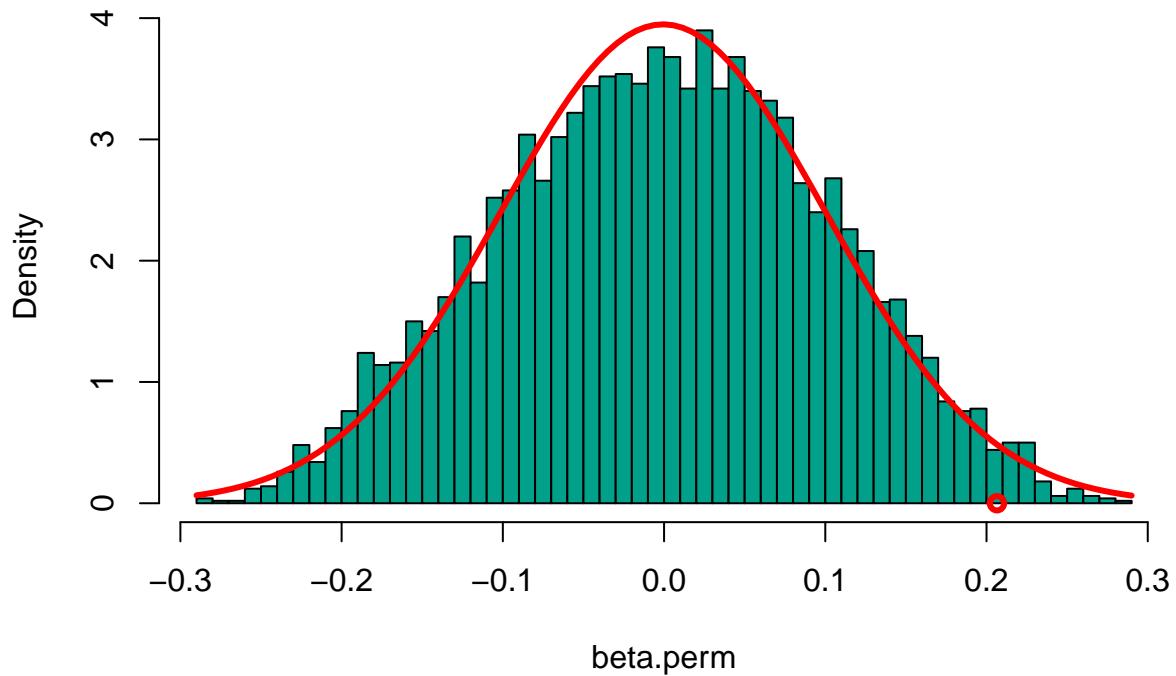
In any case, when computationally feasible, computing exact p-values is always preferable to estimating them.

## 2.4 Comparison with Parametric Linear Models

The histogram of test statistics from permuted data is well approximated by a **Gaussian** curve.

```
hist(beta.perm, 50, probability = TRUE, col = 2)
curve(dnorm(x, mean(beta.perm), sd(beta.perm)), add = TRUE, col = 1, lwd = 3)
points(beta1, 0, lwd = 3, col = 1)
```

## Histogram of beta.perm



### 2.4.1 Simple Linear Parametric Model

Assume observed values distribute around true values  $\beta_0 + \beta_1 X$  according to a Gaussian distribution:

$Y$  = linear part + normal error

$$Y = \beta_0 + \beta_1 X + \varepsilon$$

**Linear model assumptions:** -  $y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$  (true linear relationship plus error) -  $\varepsilon_i \sim N(0, \sigma^2)$  for all  $i = 1, \dots, n$  (normal errors with zero mean and constant variance/homoscedasticity)

### 2.4.2 Hypothesis Testing

If assumptions hold:

$$\hat{\beta}_1 \sim N(\beta_1, \sigma^2 / \sum(x_i - \bar{x})^2)$$

Test statistic:

$$t = \frac{\hat{\beta}_1}{\text{std.dev}(\hat{\beta}_1)} = \frac{\hat{\beta}_1}{\sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum(x_i - \bar{x})^2} / (n-2)}}$$

If  $H_0 : \beta_1 = 0$  is true,  $t \sim t(n - 2)$ .

For reaction data with  $H_1 : \beta_1 \neq 0$  (two-sided):

```
model <- lm(Reaction.Time ~ Age, data = reaction)
summary(model)
```

```
##
## Call:
```

```

## lm(formula = Reaction.Time ~ Age, data = reaction)
##
## Residuals:
##    Min     1Q Median     3Q    Max 
## -6.535 -3.364 -0.272  2.676  7.839 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 10.30135   4.04407   2.547   0.0343 *  
## Age         0.20647   0.07841   2.633   0.0300 *  
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 4.678 on 8 degrees of freedom
## Multiple R-squared:  0.4643, Adjusted R-squared:  0.3973 
## F-statistic: 6.934 on 1 and 8 DF,  p-value: 0.03003

```

Similar result, but with many more assumptions!

### 2.4.3 Permutation Test Assumptions

What model do permutation tests assume?

Under  $H_0: f(y) = f(y|x) \forall x$

Under  $H_1$ : No assumptions. For power, we hope:  $H_1: E(y|x) = \beta_0 + \beta_1 x$  with  $\beta_1 \neq 0$  for some  $x$   
i.e.,  $H_1: E(y|x) \neq E(x)E(y)$

No other assumptions about  $f(y|x)$  distribution (normality, finite moments, etc.).

## 2.5 Permutationally Equivalent Tests

```

set.seed(1)
(res_cor <- flip(Reaction.Time ~ Age, data = reaction, statTest = "cor"))

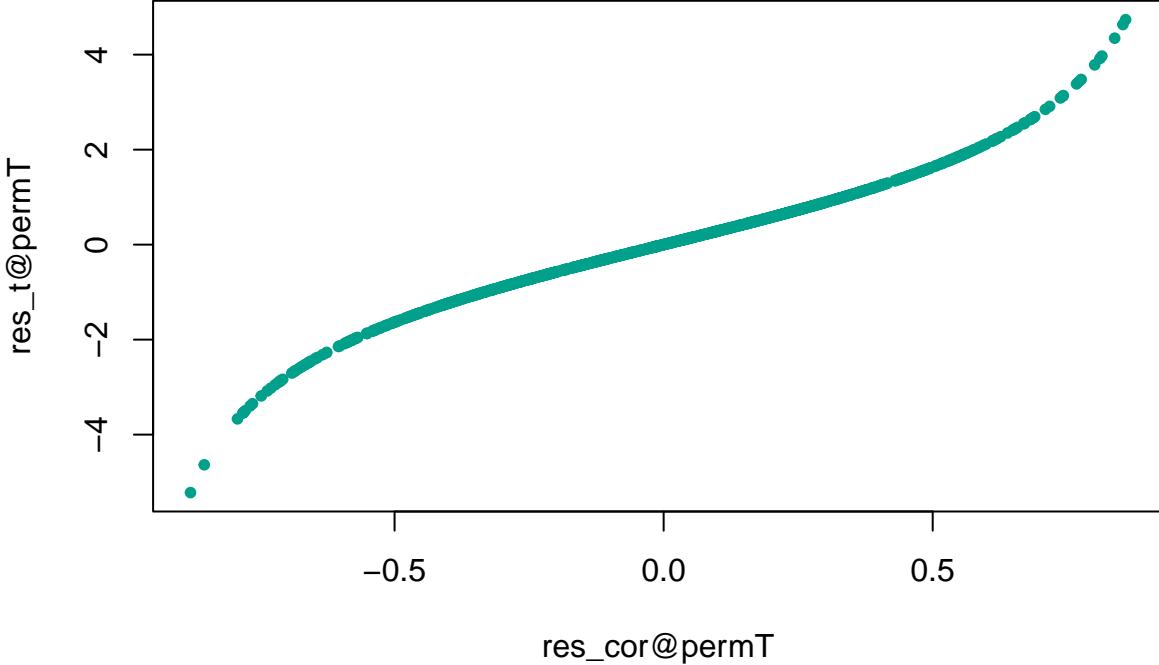
## 
##           Test  Stat tail p-value
## Reaction.Time cor 0.6814  ><  0.0410

set.seed(1)
(res_t <- flip(Reaction.Time ~ Age, data = reaction, statTest = "t"))

## 
##           Test  Stat tail p-value
## Reaction.Time   t 2.633  ><  0.0410

plot(res_cor@permT, res_t@permT, pch = 20, col = 2)

```



### 2.5.1 Conclusion

**Permutation tests:** - Different from bootstrap methods (permutation: without replacement; bootstrap: with replacement). Permutation tests have optimal properties and (usually) exact Type I error control. - General approach applicable in many contexts with minimal assumptions.

- Dedicated R packages:
  - \* **coin** <http://cran.r-project.org/web/packages/coin/index.html>
  - \* **permuco** <https://cran.r-project.org/web/packages/permuco/index.html>
  - \* **flip** <http://cran.r-project.org/web/packages/flip/index.html> (development: <https://github.com/livioivil/flip>)
  - \* **flipscores** <http://cran.r-project.org/web/packages/flipscores/index.html> (development: <https://github.com/livioivil/flipscores>)
  - \* **multcomp** <https://cran.r-project.org/web/packages/multcomp/index.html>
  - \* **GFD** <https://cran.r-project.org/web/packages/GFD/index.html>

## 3 Special Cases

### 3.1 Rank Correlation

- $n$  observations of  $y$ , interest in  $F(y|x)$ 
  - Don't need  $y_1$  and  $y_2$  to be continuous or have finite moments
- Hypotheses:
  - $H_0 : F(y|x) = F(y|x')$   $\forall x, x'$
  - $H_1 : \exists x < x' : F(y|x) < F(y|x')$  or directional alternatives
  - Test statistic: rank correlation

```

(res <- flip(Reaction.Time ~ Age, data = reaction, perms = 5000, statTest = "rank"))

##
##          Test Stat tail p-value
## Reaction.Time Wilcoxon 2.189   >< 0.0204

# Alternative using rank transformation:
(res <- flip(rank(reaction$Reaction.Time) ~ rank(reaction$Age), perms = 5000, statTest = "cor"))

##
##          Test Stat tail p-value
## rank.reaction.Reaction.Time. cor 0.7153   >< 0.0222

(cor.test(reaction$Reaction.Time, reaction$Age, method = "spearman"))

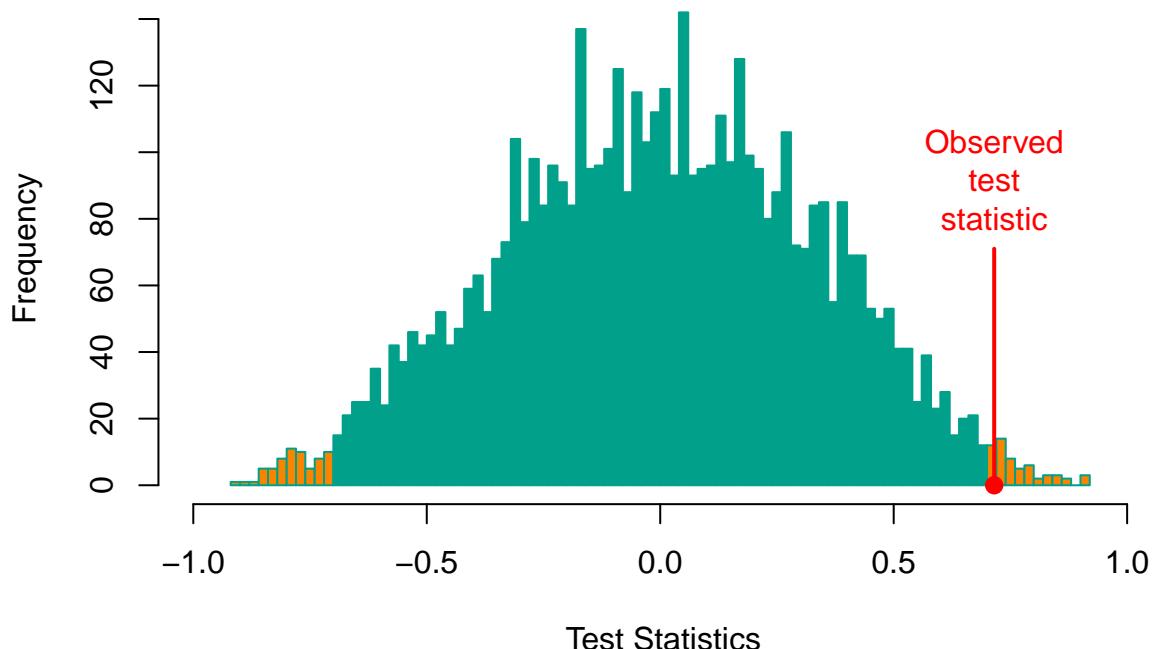
## Warning in cor.test.default(reaction$Reaction.Time, reaction$Age, method =
## "spearman"): Impossibile calcolare p-value esatti in presenza di ties

##
## Spearman's rank correlation rho
##
## data: reaction$Reaction.Time and reaction$Age
## S = 46.983, p-value = 0.02005
## alternative hypothesis: true rho is not equal to 0
## sample estimates:
##      rho
## 0.715256

plot(res)

```

### rank.reaction.Reaction.Time.



## 3.2 Two Independent Samples Problem

- Two samples:  $n_1$  observations from  $y_1$ ,  $n_2$  from  $y_2$ 
  - No continuity or finite moment requirements
- Hypotheses:
  - $H_0 : F(y_1) = F(y_2)$
  - $H_1 : F(y_1) \neq F(y_2)$  (or directional)
- Test statistics:
  - Standardized mean difference (t-statistic)
- Estimated slope coefficient (group labels as dummy predictor)
- Other permutationally equivalent statistics

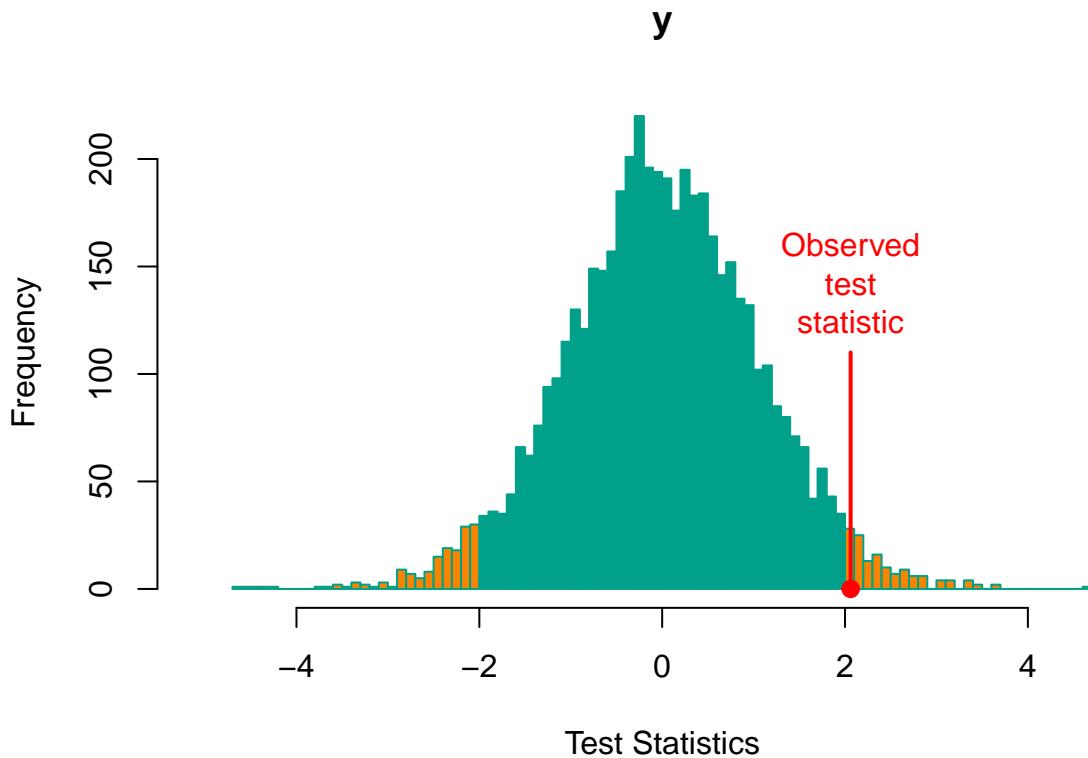
```
data("seeds")
seeds <- na.omit(seeds)
(res <- flip(y ~ grp, data = seeds, perms = 5000))

##
##  Test Stat tail p-value
## y      t 2.061    >< 0.0526

summary(lm(y ~ grp, data = seeds))

##
## Call:
## lm(formula = y ~ grp, data = seeds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.331 -2.931 -1.651  4.663  7.863
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.147     1.242   8.168   9e-09 ***
## grp          3.345     1.623   2.061    0.049 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.303 on 27 degrees of freedom
## Multiple R-squared:  0.136, Adjusted R-squared:  0.104
## F-statistic: 4.249 on 1 and 27 DF,  p-value: 0.04903

plot(res)
```



### 3.2.1 Rank Test

Can use rank-based statistics? Yes—equivalent to rank tests but with exact distribution (no tie limitations).

```
(res <- flip(y ~ grp, data = seeds, statTest = "rank", perms = 5000))

##
##      Test Stat tail p-value
## y Wilcoxon 2.148   < 0.0292

wilcox.test(y ~ grp, data = seeds)

## Warning in wilcox.test.default(x = DATA[[1L]], y = DATA[[2L]], ...): non è
## possibile calcolare p-value esatto in presenza di ties

##
## Wilcoxon rank sum test with continuity correction
##
## data: y by grp
## W = 53.5, p-value = 0.03353
## alternative hypothesis: true location shift is not equal to 0
```

## 3.3 Chi-square and Other Categorical Methods

```

data("seeds")
seeds$Germinated <- !is.na(seeds$x)
seeds$Germinated <- factor(seeds$Germinated)
seeds$grp <- factor(seeds$grp)
table(seeds$grp, seeds$Germinated)

##          FALSE TRUE
## 0      8    12
## 1      3    17

chisq.test(seeds$grp, seeds$Germinated)

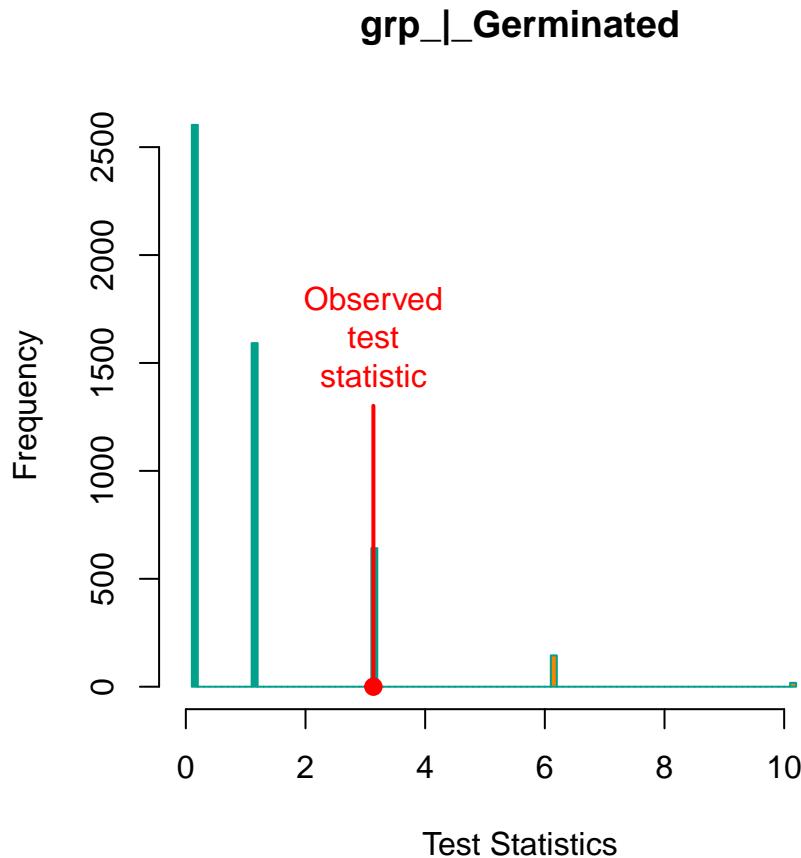
##
##  Pearson's Chi-squared test with Yates' continuity correction
##
## data:  seeds$grp and seeds$Germinated
## X-squared = 2.0063, df = 1, p-value = 0.1567

(res <- flip(Germinated ~ grp, data = seeds, statTest = "Chisq", perms = 5000))

##
##                                Test  Stat tail p-value
## grp_|_Germinated Chi Squared 3.135    > 0.1610

plot(res)

```



...and Fisher's exact test:

```
fisher.test(seeds$grp, seeds$Germinated)$p.value

## [1] 0.1551874

flip(Germinated ~ grp, data = seeds, perms = 5000)

## 
##          Test   Stat tail p-value
## GerminatedFALSE    t -1.798   <  0.1596
## GerminatedTRUE      t  1.798   <  0.1596
```

### 3.4 ANOVA (C-sample)

Example: 3 Age groups: young [18 – 35), middle [35 – 60), old [60 – 100)

- C samples:  $n_i$  observations from  $y_i$  ( $i = 1, \dots, C$ )
  - No continuity or finite moment requirements
- Hypotheses:
  - $H_0 : F(y_i) = F(y_j) \forall (i, j)$

- $H_1 : \exists(i, j) : F(y_i) \neq F(y_j)$

- Test statistics:

- F-statistic
- $R^2$
- Other permutationally equivalent statistics
- Rank-based alternatives

```
reaction$AgeCateg <- cut(reaction$Age, c(18, 35, 65, 100), right = FALSE)
(res <- flip(Reaction.Time ~ AgeCateg, data = reaction, perms = 5000, statTest = "ANOVA"))
```

```
##                                     Test Stat tail p-value
## Reaction.Time      F 4.02      > 0.0780
summary(lm(Reaction.Time ~ AgeCateg, data = reaction))
```

```
##
## Call:
## lm(formula = Reaction.Time ~ AgeCateg, data = reaction)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -6.495 -3.279  0.465  2.246  6.112
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 16.157     2.331   6.932 0.000225 ***
## AgeCateg[35,65] 4.428     3.296   1.343 0.221144
## AgeCateg[65,100] 11.418     4.037   2.828 0.025478 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.662 on 7 degrees of freedom
## Multiple R-squared:  0.5346, Adjusted R-squared:  0.4016
## F-statistic:  4.02 on 2 and 7 DF,  p-value: 0.06878
```

### 3.4.1 Stochastic Ordering

- Same assumptions as ANOVA
- Hypotheses:
  - Same  $H_0 : F(y_i) = F(y_j) \forall(i, j)$
  - But  $H_1 : \exists(i, j) : F(y_i) < F(y_j)$  (or  $>$ )

(More details on NPC later)

```
(res <- flip(Reaction.Time ~ AgeCateg, data = reaction, perms = 5000, tail = 1))
```

```
##                                     Test Stat tail p-value
## Reaction.Time_ | _AgeCateg.[35,65].      t 0.1423      > 0.4336
## Reaction.Time_ | _AgeCateg.[65,100].      t 2.2444      > 0.0220
```

```

npc(res)

##          comb.funct nVar  Stat p-value
## V1      Fisher     2 4.652  0.0202

```

### 3.5 Stratified Permutations (Discrete Nuisances)

Test  $X = \text{Age}$  with  $Z = \text{Gender}$  as nuisance in `reaction` data.

Under  $H_0$ :  $f(y|x, z) = f(y|x', z) = f(y|z) \forall (x, x')$

Thus, under  $H_0$ ,  $f(y_i) = f(y_j)$  only if  $z_i = z_j$  (same gender).

Can we permute as before? NO. Permute only within strata defined by  $Z$ .

**Remark:** - No linear nuisance effect assumption - Allow heteroscedastic errors across strata

Test statistic remains unchanged.

```
(res <- flip(Reaction.Time ~ Age, Strata = ~Gender, data = reaction, perms = 5000))
```

```

##          Test   Stat tail p-value
## Reaction.Time    t  2.633  ><  0.0718

```

Alternative model (more on NPC later):

```
(res <- flip(Reaction.Time ~ Age*Gender, Strata = ~Gender, data = reaction, perms = 5000))
```

```

##          Test   Stat tail p-value
## Reaction.Time | _Age        t  2.4826  ><  0.0750
## Reaction.Time | _Age:Gender.M. t -0.6518  ><  0.3394

```

```
npc(res)
```

```

##          comb.funct nVar  Stat p-value
## V1      Fisher     2 3.671  0.1406

```

## 4 Paired samples

Understanding the Model:  $Y = \gamma_0 + \gamma_1 Z + \beta X + \varepsilon$

Consider the model: -  $Z$  represents Subject ID ( $i = 1, \dots, n$ ) -  $X$  indicates pre/post-treatment status (typically coded as 0 for pre-treatment, 1 for post-treatment) - Each subject in  $Z$  has one pre-treatment and one post-treatment observation

The null hypothesis to test is  $H_0 : \beta = 0$ .

When permutations are constrained within levels of  $Z$  (within each subject), this becomes equivalent to flipping the sign of the difference between post- and pre-treatment measurements.

Let  $D$  be the vector of these differences with  $n$  observations. The null hypothesis can be rewritten as  $H_0 : E(D) = 0$ .

## 4.1 Advantages of this approach

One major advantage is that we don't need to estimate the Fisher Information (i.e., the residual variance). Let's demonstrate this with a paired t-test example:

```

n <- 20
y <- rnorm(n)

# Generate sign flips for permutation testing
FLIPS <- flipscores:::make_flips(n, 1000)
head(FLIPS)

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
## [1,]    1    1    1    1    1    1    1    1    1    1    1    1    1    1    1
## [2,]   -1    1   -1   -1    1    1   -1    1   -1    1    1    1   -1    1
## [3,]   -1   -1    1    1    1   -1   -1    1   -1    1   -1   -1    1    1
## [4,]    1   -1   -1    1    1    1    1   -1   -1    1    1   -1   -1    1    1
## [5,]    1   -1    1   -1   -1    1   -1   -1    1    1   -1    1   -1   -1
## [6,]    1    1    1   -1    1   -1   -1   -1    1    1    1   -1    1   -1
##     [,15] [,16] [,17] [,18] [,19] [,20]
## [1,]    1    1    1    1    1
## [2,]   -1   -1    1   -1   -1
## [3,]   -1    1    1    1   -1    1
## [4,]   -1    1    1    1    1
## [5,]    1    1   -1    1    1
## [6,]   -1    1    1   -1   -1

# Calculate test statistics for each permutation
Tstat <- FLIPS %*% y
flipscores:::t2p(Tstat)

## [1] 0.124

# Monte Carlo simulation to check Type I error rate
MCMC <- 10000
Y <- matrix(rnorm(n * MCMC), n, MCMC)
Tstat <- FLIPS %*% Y
p.values <- apply(Tstat, 2, flipscores:::t2p)
mean(p.values < 0.05)

## [1] 0.047

# Lightweight t-test function
t.test.light <- function(Y, tail = 1) {
  sd <- apply(Y, 2, sd, na.rm = TRUE)
  n <- nrow(Y)
  ts <- colMeans(Y, na.rm = TRUE) / sd * sqrt(n)
  pt(-ts, df = n - 1)
}

# Test with heteroscedastic data
Y <- replicate(MCMC, rnorm(n, sd = exp(1:n)))

```

```

Tstat <- FLIPS %*% Y
p.values <- apply(Tstat, 2, flipscores:::t2p)
p.values_param <- t.test.light(Y)

# Compare Type I error rates
mean(p.values < 0.05)

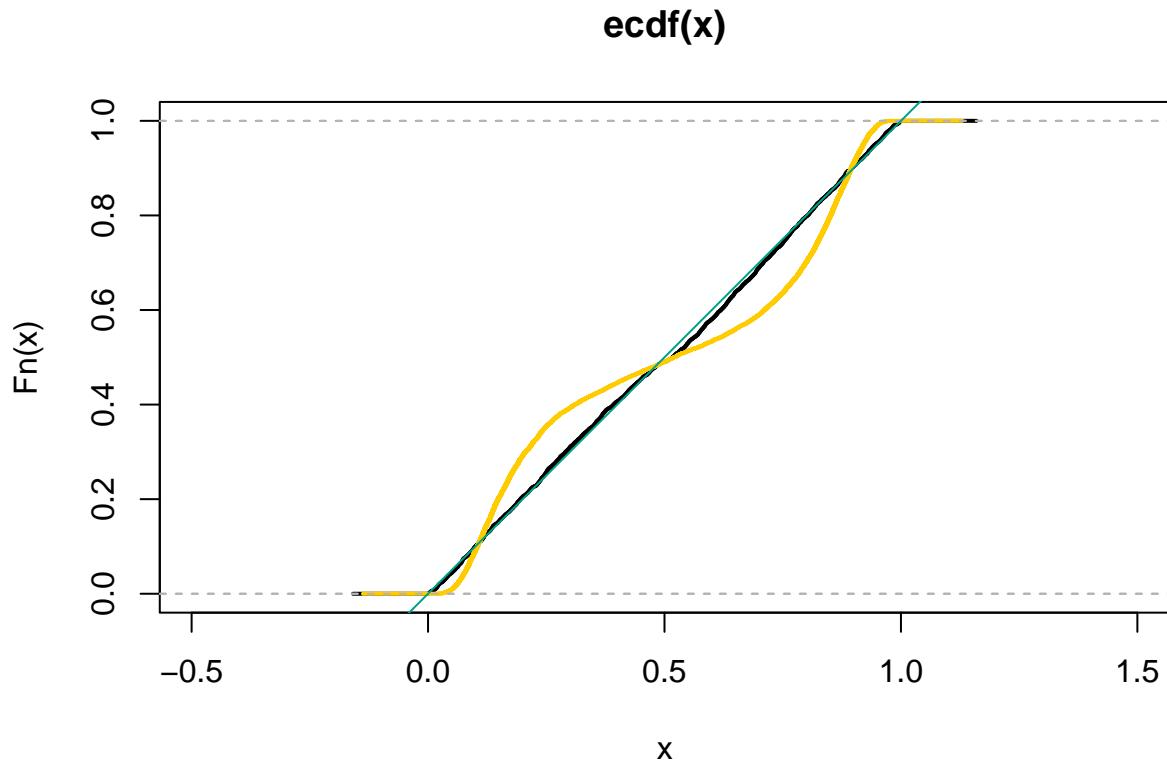
## [1] 0.0438

mean(p.values_param < 0.05)

## [1] 0.0095

# Visualize p-value distributions
plot.ecdf(p.values,lwd=2,asp=1)
plot.ecdf(p.values_param,lwd=2, add = TRUE, col = 3)
abline(0, 1, col = 2)

```



In this example, we compare a permutation-based approach (sign-flipping test) with a parametric t-test:

- 1. Permutation Test:** We generate sign flips that correspond to within-subject permutations. The test statistic is calculated for each sign-flipped version of the data, creating a null distribution against which we compare our observed statistic.
- 2. Parametric t-test:** We use a standard t-test that assumes normality and constant variance.

3. **Heteroscedasticity Simulation:** The second part of the code generates data with increasing variance across observations (`sd = exp(1:n)`). This violates the constant variance assumption of the parametric t-test.

The empirical cumulative distribution function (ECDF) plot shows how the p-values from both methods compare. Under the null hypothesis with correct assumptions, both should produce uniformly distributed p-values (follow the diagonal line). When assumptions are violated, the parametric test may produce inflated Type I error rates, while the permutation test maintains correct error control due to its distribution-free nature.

**Key Insights:** - The permutation test doesn't require estimating residual variance - It remains valid even when parametric assumptions (like homoscedasticity) are violated - The method is particularly powerful for paired data where within-subject comparisons are natural - Sign-flipping is equivalent to permuting within subjects when testing for treatment effects in paired designs

## 4.2 Repeated measures and mixed models

More properly this approach is known as Random coefficient Analysis / Group-level analysis

- Basso & Finos (2012). Exact Multivariate Permutation Tests for Fixed Effects in Mixed-Models. *Communications in Statistics - Theory and Methods*, 41: 2991 - 3001.
- Finos & Basso (2014) Permutation tests for between-unit fixed effects in multivariate generalized linear mixed models. *Stat Comput* 24, 941–952.

A more complete approach:

- Andreella, Goeman, Hemerik, Finos (2025) Robust Inference for Generalized Linear Mixed Models: A “Two-Stage Summary Statistics” Approach Based on Score Sign Flipping. *Psychometrika*, 1-23.

## 5 Multivariate Testing

### 5.1 Seeds Data

```
# install.packages("flip")
library(flip)
```

Remove NAs:

```
data(seeds, package = "flip")
seeds <- na.omit(seeds)
seeds
```

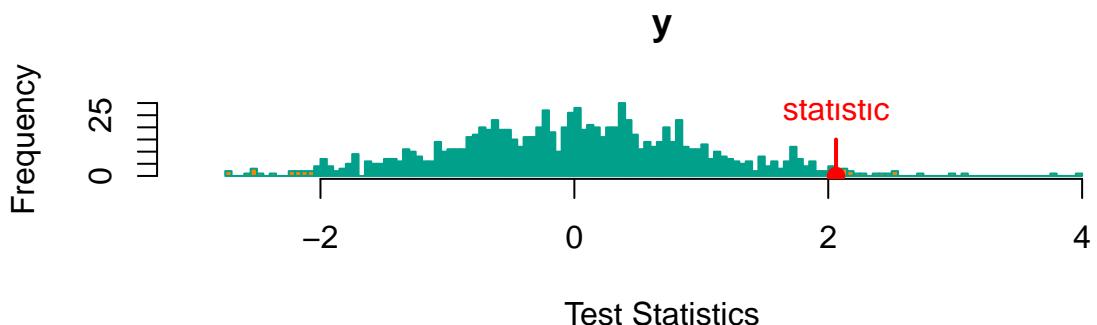
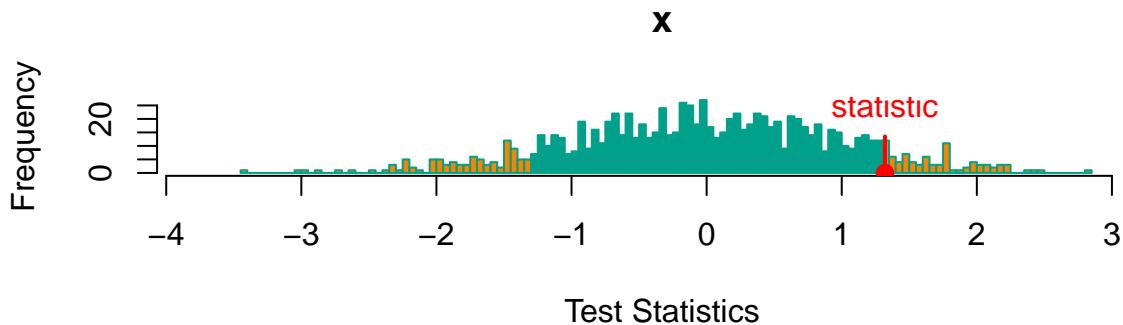
```
##      grp      x      y
## 9      0 6.03 12.54
## 10     0 4.20 14.81
## 11     0 4.49 16.71
## 12     0 2.00  7.53
## 13     0 2.84  7.02
```

```
## 14 0 3.88 8.09
## 15 0 2.04 5.76
## 16 0 5.48 18.01
## 17 0 2.31 8.81
## 18 0 1.90 8.17
## 19 0 1.75 6.62
## 20 0 3.02 7.69
## 24 1 3.31 18.49
## 25 1 6.56 19.20
## 26 1 3.16 9.85
## 27 1 4.07 15.83
## 28 1 2.09 6.16
## 29 1 6.72 17.58
## 30 1 3.93 19.29
## 31 1 2.56 10.77
## 32 1 8.30 18.31
## 33 1 4.21 10.56
## 34 1 1.86 9.48
## 35 1 3.09 12.54
## 36 1 5.09 18.35
## 37 1 4.08 11.84
## 38 1 3.63 11.44
## 39 1 2.61 7.66
## 40 1 5.21 12.00
```

## 5.2 Marginal vs Joint Distribution

Use permutation methods to test for group differences (grp) on both x and y:

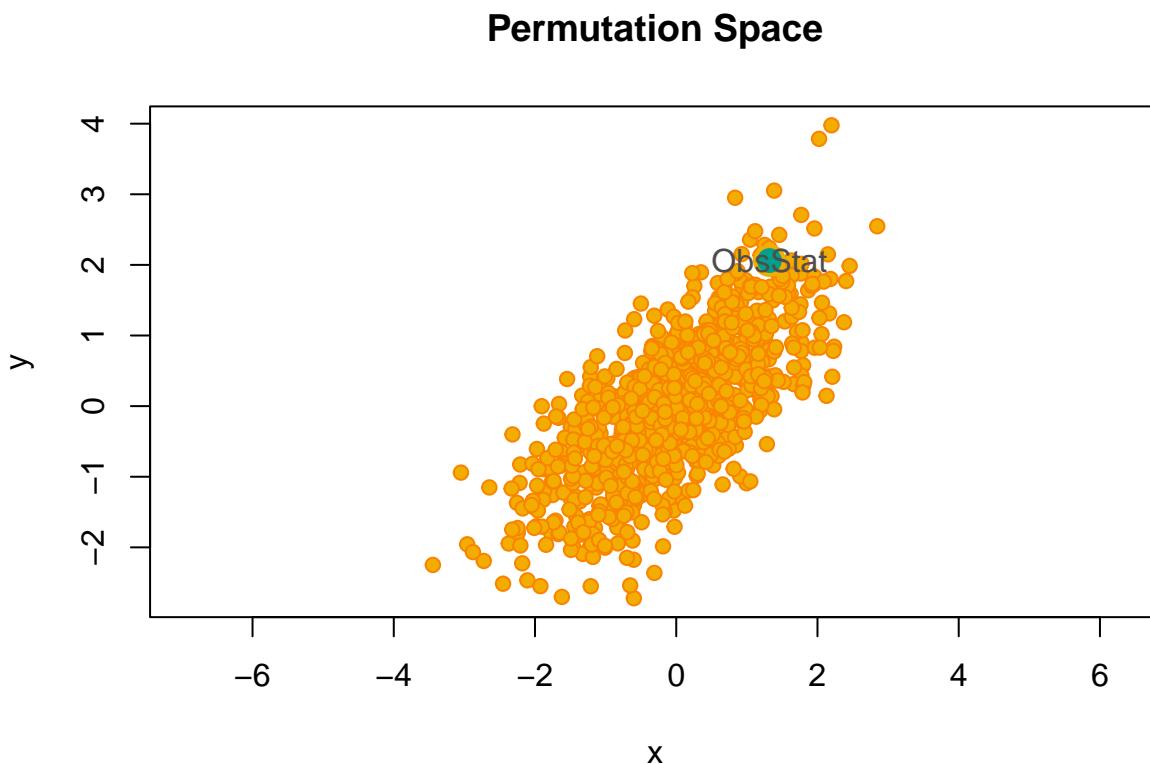
```
library(flip)
res <- flip(. ~ grp, data = seeds, flipReturn = list(permP = TRUE, permT = TRUE))
hist(res)
```



```
# flipReturn = list(permP = TRUE, permT = TRUE) not strictly needed; useful later
```

We can test the two variables separately but lack an overall p-value (is there ANY difference?).

```
plot(res)
```



Next we'll see:

- How to combine p-values (e.g., Fisher's combining function) for global hypothesis testing
- How to use closed testing procedures to adjust p-values: which variables differ?

### 5.3 Rejection Regions (and Overall Testing)

In univariate settings, defining ‘far from null’ (usually from test statistic = 0) is straightforward. In multivariate settings, there are multiple (no uniformly best) approaches.

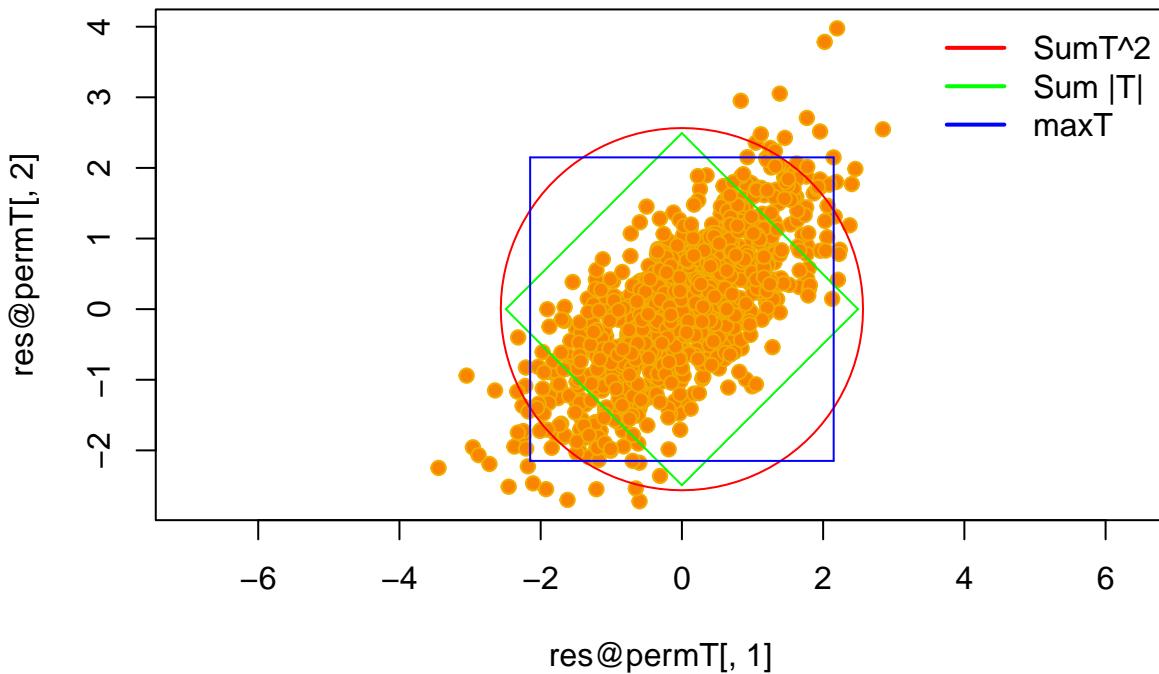
```
# install.packages("plotrix")
library("plotrix")

## Warning: il pacchetto 'plotrix' è stato creato con R versione 4.5.2

res.sumt2 <- npc(res, "sumT2", flipReturn = list(permP = TRUE, permT = TRUE))
limsumt2 <- res.sumt2@permT[which.min(abs(res.sumt2@permP - 0.05))]
res.sumt <- npc(res, "sumT", flipReturn = list(permP = TRUE, permT = TRUE))
limsumt <- res.sumt@permT[which.min(abs(res.sumt@permP - 0.05))]
res.maxt <- npc(res, "maxT", flipReturn = list(permP = TRUE, permT = TRUE))
limmaxt <- res.maxt@permT[which.min(abs(res.maxt@permP - 0.05))]

plot(res@permT[,1], res@permT[,2], col = "#F2AD00", bg = "#F98400", pch = 21,
     main = "Some Rejection Regions (alpha = .05)", asp = 1)
draw.circle(0, 0, limsumt2^.5, border = "red")
segments(c(limsumt, -limsumt, limsumt, -limsumt), c(0, 0, 0, 0),
         c(0, 0, 0, 0), c(limsumt, -limsumt, -limsumt, limsumt), col = "green")
segments(c(limmaxt, -limmaxt, -limmaxt, limmaxt), c(limmaxt, limmaxt, -limmaxt, -limmaxt),
         c(-limmaxt, -limmaxt, limmaxt, limmaxt), c(limmaxt, -limmaxt, -limmaxt, limmaxt), col = "blue")
legend("topright", legend = c("SumT^2", "Sum |T|", "maxT"),
       col = c("red", "green", "blue"), bty = "n", lwd = 2)
```

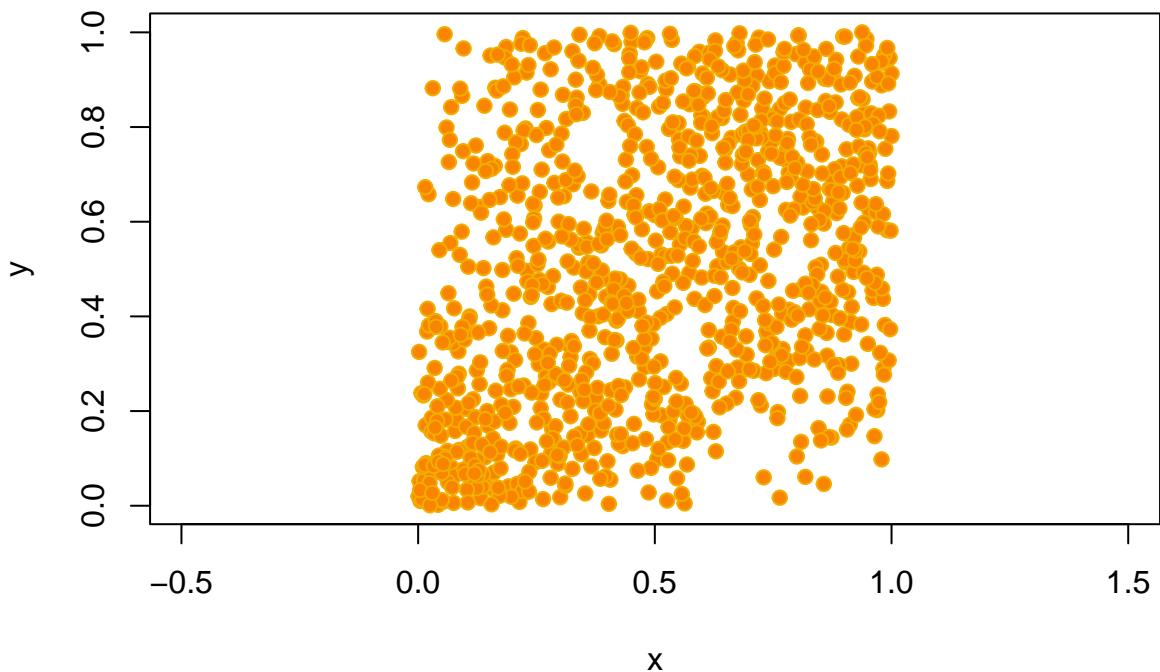
## Some Rejection Regions (alpha = .05)



**REMARK:** We can derive the p-value distribution by computing p-values for each test statistic (observed and permuted data). This yields the multivariate p-value distribution:

```
plot(res@permP, col = "#F2AD00", bg = "#F98400", pch = 21,
     main = "Joint Distribution of P-values", asp = 1)
```

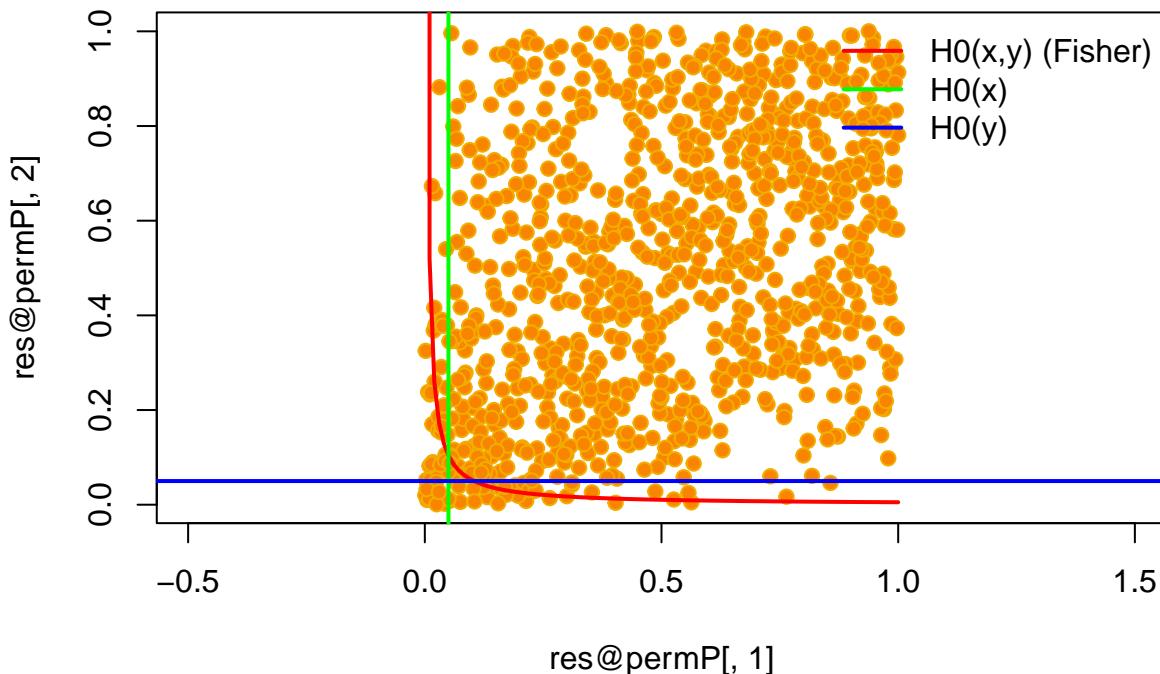
## Joint Distribution of P-values



### 5.3.1 Fisher Combining Function

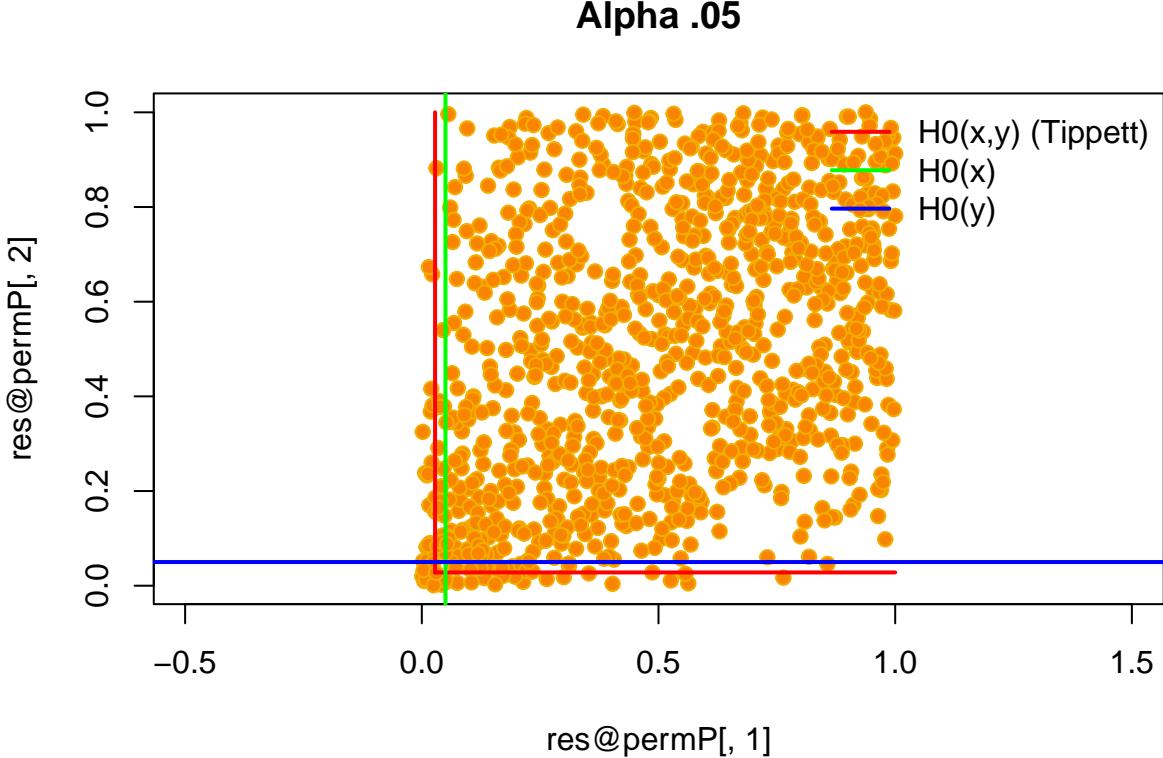
Examine rejection regions for univariate tests and Fisher combination. Intersection of each univariate test with Fisher region defines closed testing rejection region (adjusted for multiple testing).

**Alpha .05**



### 5.3.2 Tippett (min-p) Combining Function

Examine rejection regions for univariate tests and Tippett combination. Intersection defines closed testing rejection region. This coincides with Westfall & Young shortcut.



## 6 FWER Control via Permutation Tests

### 6.1 Permutation Bonferroni

Bonferroni is conservative:

- **Bonferroni bound:** Reject if p-value  $\leq \alpha/m$
- **By Boole's inequality:** Guaranteed FWER  $\leq \alpha$ , but often FWER  $< \alpha$
- **Can we improve?:** Reject if p-value  $\leq \tilde{\alpha} > \alpha/m$  while maintaining FWER control
- **Yes:** Via permutations

### 6.2 Improved Bonferroni

- **Reduced  $\alpha$ :** Reject  $H_i$  if  $p_i \leq \tilde{\alpha}$
- **FWER control?:**

$$\text{FWER} = P(p_i \leq \tilde{\alpha} \text{ for at least one true } H_i)$$

$$= P\left(\bigcup_{i \in T} \{p_i \leq \tilde{\alpha}\}\right)$$

$$= P\left(\min_{i \in T} p_i \leq \tilde{\alpha}\right) \leq \alpha$$

- **How to determine  $\tilde{\alpha}$ ?**: Use permutations to find minimum p-value distribution

### 6.3 Multiple Testing via Permutations

**Single-step min-P method:**

1. Calculate smallest p-value  $m$  for real data
2. Randomly permute data
3. Calculate new p-values for all tests on permuted data
4. Calculate smallest p-value  $m^\pi$  for permuted data
5. Repeat permutations many times (e.g.,  $k = 1000$ ):  $m_1^\pi, \dots, m_k^\pi$
6. Calculate  $\tilde{\alpha}$  as  $\alpha$ -quantile of  $m_1^\pi, \dots, m_k^\pi$

**Multiple testing result:** Reject all hypotheses with (non-permuted) p-values  $\leq \tilde{\alpha}$

### 6.4 P-value Correlation Structure

**Permutation:**

- Destroys covariate-response correlation
- Preserves covariate correlations

**Consequence:**

- P-values of correlated tests remain correlated in permutations
- Minimum p-value distribution correctly accounts for correlations

**When is improvement over Bonferroni large?:**

- Negatively correlated p-values: typically no gain
- Independent p-values: minimal gain
- Positively correlated p-values: potentially large gain

### 6.5 Improved Holm: Westfall & Young

*Westfall PH, Young SS (1993) Resampling-Based Multiple Testing: Examples and Methods for p-Value Adjustment. Wiley Sequential permutation multiple testing:*

- **Single-step**: Permutation equivalent of Bonferroni

- **Holm equivalent:** Westfall & Young method
- **min-P algorithm:**

1. Start with all hypotheses
2. Repeat:
  - Perform single-step min-P to calculate  $\tilde{\alpha}$
  - Reject hypotheses with p-value  $\leq \tilde{\alpha}$
  - Remove rejected hypotheses
3. Until no new rejections

## 6.6 General Framework: Closed Testing

*Marcus R, Peritz E, Gabriel KR (1976). On closed testing procedures with special reference to ordered analysis of variance. Biometrika 63: 655-660.*

Test each node with any multivariate permutation test.

Westfall & Young is a special case of closed testing (each node uses min-p/Tippett or max-T combining function).

### 6.6.1 Closure Set

ABC

AB

AC

BC

A

B

C

Adjusted  $\tilde{p}_A = \max(p_A, p_{AB}, p_{AC}, p_{ABC})$

In our data:

```
(res <- flip.adjust(res, method = "Fisher"))
```

```
##  
##   Test  Stat tail p-value Adjust:Fisher  
## x     t 1.320  ><  0.1710      0.1710  
## y     t 2.061  ><  0.0350      0.0620
```

```
(res <- flip.adjust(res, method = "maxT"))
```

```
##  
##   Test  Stat tail p-value Adjust:Fisher Adjust:maxT  
## x     t 1.320  ><  0.1710      0.1710      0.1710  
## y     t 2.061  ><  0.0350      0.0620      0.0600
```

Conclusion

**Accounting for dependencies:** Adjusted p-values become lower (more rejections).

**When?:**

- Negative correlation: generally no gain
- Independent p-values: little or no gain
- Positive correlation: substantial gain (note: two-sided tests with negatively correlated test statistics yield positively correlated p-values)

**Real data:** Often correlated variables → permutations advantageous

**How?:** R: `library(flip); flip(); flip.adjust()`

## 7 Case Study: Pharmacokinetic Study of Carbidopa

Description:

<http://webserv.jcu.edu/math//faculty/TShort/Bradstreet/part2/part2-table6.html>

12 healthy male subjects in three-period crossover design receiving three graded doses (25, 50, 100 mg) of Carbidopa q8h. Seven-day washout between periods. Pharmacokinetic variables AUC, Cmax, Tmax calculated from plasma concentrations at 0, 0.5, 1, 1.5, 2, 3, 4, 5, 6, 7, 8 hours postdosing after second dose on day 6.

Dataset:

<http://webserv.jcu.edu/math//faculty/TShort/Bradstreet/part2/Bradp2t6.txt>

Analyze without accounting for study periods (randomized within subjects).

Research questions: 1. Is there a dose response for AUC, Cmax, or Tmax? Overall? 2. Can dose proportionality be established? (Fit linear model for each endpoint, discuss results)

## 7.1 Solution

Address both questions with single analysis: linear model (accounting for individual variability) on log-transformed endpoints.

```
# Read and prepare data
dati <- read.table("http://webserv.jcu.edu/math//faculty/TShort/Bradstreet/part2/Bradp2t6.txt",
                    skip = 1, header = TRUE)
dati <- cbind(dati[,1], matrix(as.matrix(dati[,-1]), nrow(dati)*3, 4))
colnames(dati) <- c("Sub", "Dose", "AUC", "Cmax", "Tmax")
dati <- as.data.frame(dati)
str(dati)

## 'data.frame':   36 obs. of  5 variables:
## $ Sub : num  1 2 3 4 5 6 7 8 9 10 ...
## $ Dose: num  100 25 50 50 50 25 100 25 50 25 ...
## $ AUC : num  604 140 386 175 605 ...
## $ Cmax: num  137 44.4 86.6 46.4 194 44.9 318 29 119 58.4 ...
## $ Tmax: num  1.5 1 1.5 1.5 0.5 1 1 1 2 2 ...

# Log-transform responses (linear relationship indicates proportionality)
dati[,3:5] <- log(dati[,3:5])

# Descriptives and plots
summary(dati[,-1])

##          Dose            AUC            Cmax            Tmax
## Min.   : 25.00    Min.   :4.337    Min.   :3.219    Min.   :-0.6931
## 1st Qu.: 25.00   1st Qu.:5.156    1st Qu.:3.966    1st Qu.: 0.0000
## Median : 50.00   Median :5.886    Median :4.485    Median : 0.2027
## Mean   : 58.33   Mean   :5.873    Mean   :4.547    Mean   : 0.2474
## 3rd Qu.:100.00   3rd Qu.:6.539    3rd Qu.:5.280    3rd Qu.: 0.6931
## Max.   :100.00   Max.   :7.335    Max.   :5.989    Max.   : 1.0986

by(dati[,3:5], dati$Dose, summary)

## dati$Dose: 25
##          AUC            Cmax            Tmax
## Min.   :4.337    Min.   :3.219    Min.   :-0.6931
## 1st Qu.:4.803    1st Qu.:3.390    1st Qu.: 0.0000
## Median :4.972    Median :3.801    Median : 0.0000
## Mean   :5.051    Mean   :3.783    Mean   : 0.2071
## 3rd Qu.:5.289    3rd Qu.:4.022    3rd Qu.: 0.6931
## Max.   :5.818    Max.   :4.464    Max.   : 0.6931
## -----
## dati$Dose: 50
##          AUC            Cmax            Tmax
## Min.   :5.133    Min.   :3.837    Min.   :-0.6931
## 1st Qu.:5.670    1st Qu.:4.374    1st Qu.: 0.0000
## Median :5.886    Median :4.484    Median : 0.2027
## Mean   :5.815    Mean   :4.479    Mean   : 0.1689
## 3rd Qu.:5.967    3rd Qu.:4.625    3rd Qu.: 0.4055
```

```

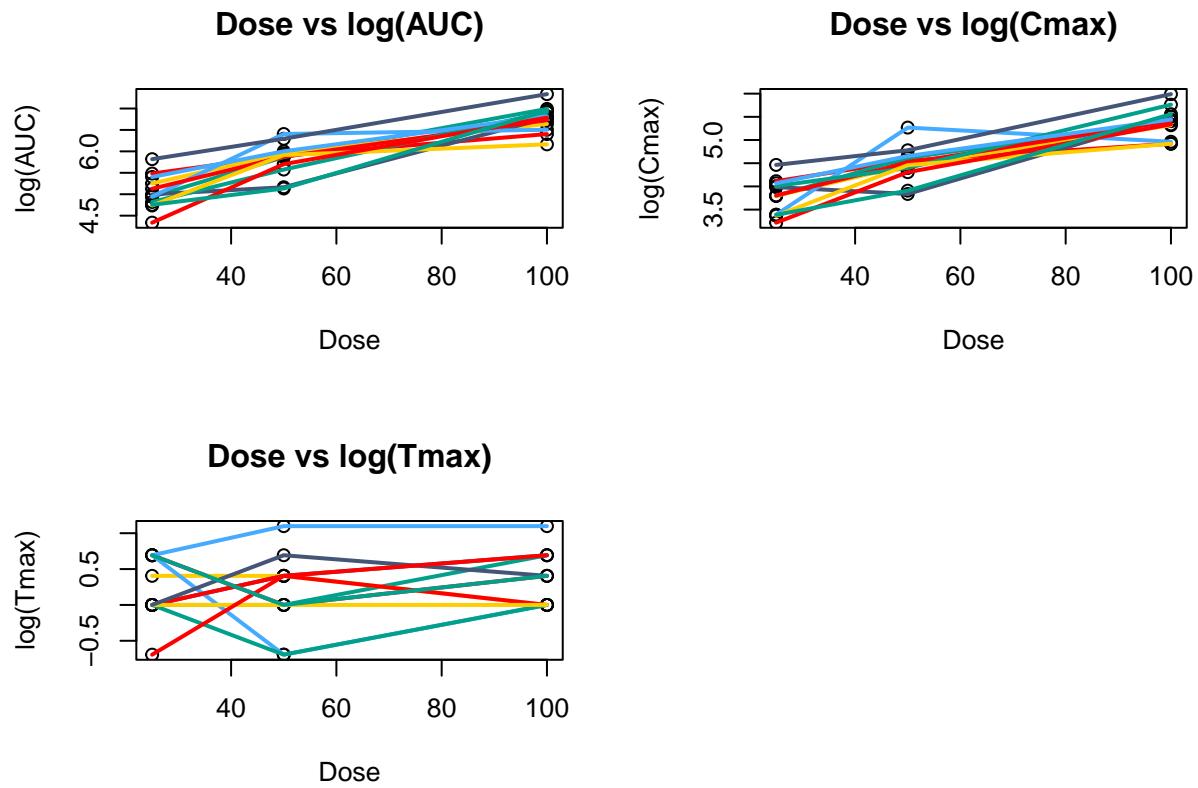
##  Max.    :6.405   Max.    :5.268   Max.    : 1.0986
## -----
##  dati$Dose: 100
##          AUC           Cmax           Tmax
##  Min.    :6.164   Min.    :4.920   Min.    :0.0000
##  1st Qu.:6.607   1st Qu.:5.229   1st Qu.:0.0000
##  Median  :6.782   Median  :5.412   Median  :0.4055
##  Mean    :6.751   Mean    :5.378   Mean    :0.3662
##  3rd Qu.:6.922   3rd Qu.:5.515   3rd Qu.:0.6931
##  Max.    :7.335   Max.    :5.989   Max.    :1.0986

par(mfrow = c(2,2))
plot(dati$Dose, dati$AUC, ylab = "log(AUC)", xlab = "Dose", main = "Dose vs log(AUC)")
temp=sapply(unique(dati$Sub), function(s) {
d <- subset(dati, Sub == s)
d <- d[order(d$Dose),]
lines(d$Dose, d$AUC, col = s, lwd = 2)
})

plot(dati$Dose, dati$Cmax, ylab = "log(Cmax)", xlab = "Dose", main = "Dose vs log(Cmax)")
temp=sapply(unique(dati$Sub), function(s) {
d <- subset(dati, Sub == s)
d <- d[order(d$Dose),]
lines(d$Dose, d$Cmax, col = s, lwd = 2)
})

plot(dati$Dose, dati$Tmax, ylab = "log(Tmax)", xlab = "Dose", main = "Dose vs log(Tmax)")
temp=sapply(unique(dati$Sub), function(s) {
d <- subset(dati, Sub == s)
d <- d[order(d$Dose),]
lines(d$Dose, d$Tmax, col = s, lwd = 2)
})

```

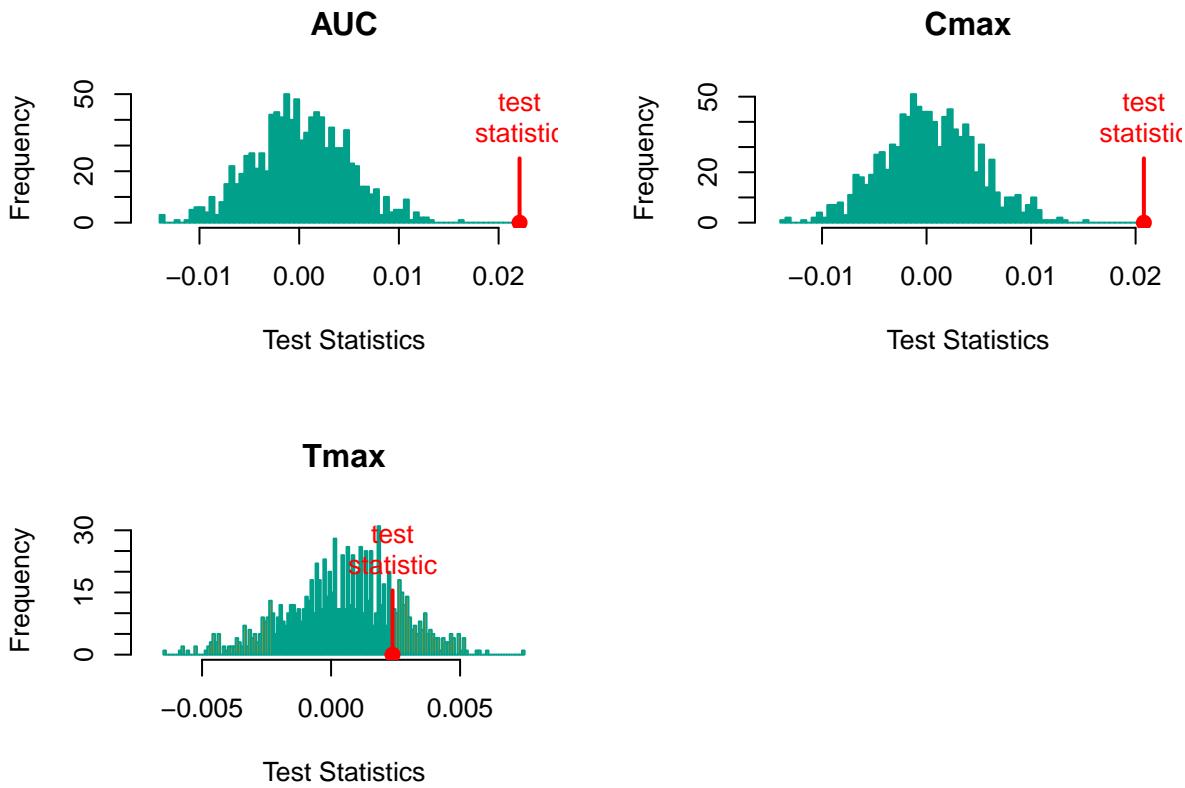


Simple solution:

```
library(flip)
res <- flip(. ~ Dose, data = dati, Strata = ~Sub, statTest = "coeff")
summary(res)

## Call:
##  flip(Y = . ~ Dose, data = dati, statTest = "coeff", Strata = ~Sub)
## 999 permutations.
##
##      Test   Stat tail p-value sig.
## AUC coeff 0.0221  >< 0.0010 *** 
## Cmax coeff 0.0208  >< 0.0010 *** 
## Tmax coeff 0.0024  >< 0.2950

# statTest = "coeff": estimated linear model coefficient
hist(res)
```



Multivariate: - Overall

```
res <- flip.adjust(res)
npc(res, "Fisher")
```

```
##
##      comb.funct nVar   Stat p-value
## V1       Fisher     3 15.04  0.0010
```

Dose effect exists overall.

- By endpoint (closed testing with max-t). Try different methods (e.g., `method = "Fisher"`) and compare `method = "minP"` with `method = "Holm"`.

```
res <- flip.adjust(res, method = "holm")
res <- flip.adjust(res, method = "Fisher")
summary(res)
```

```
## Call:
## flip(Y = . ~ Dose, data = dati, statTest = "coeff", Strata = ~Sub)
## 999 permutations.
##
##      Test   Stat tail p-value Adjust:maxT Adjust:holm Adjust:Fisher sig.
## AUC  coeff 0.0221    ><  0.0010      0.0010      0.0030      0.0010  ***
## Cmax coeff 0.0208    ><  0.0010      0.0010      0.0030      0.0010  ***
## Tmax coeff 0.0024    ><  0.2950      0.2950      0.2950      0.2950
```

AUC and Cmax show significant effects after multiplicity correction; Tmax does not.

## 8 Minimal Bibliography

### **Grounding Theory:**

- Pesarin (2001) *Multivariate Permutation Tests: With Applications in Biostatistics*. Wiley, New York

### **Alternative permutation testing approach:**

- Hemerik J, Goeman J. Exact testing with random permutations. *Test*. 2018;27(4):811-825. doi:10.1007/s11749-017-0571-1

### **Flexible GLM approach via sign-flip score test:**

- Hemerik, Goeman and Finos (2020) Robust testing in generalized linear models by sign flipping score contributions. *Journal of the Royal Statistical Society Series B* 82(3). DOI: 10.1111/rssb.12369

Implemented in R package **flipscores**:

<https://cran.r-project.org/web/packages/flipscores/index.html>

Development version: <https://github.com/livioivil/flipscores>

### **Permutation regression review:**

- Winkler AM, Ridgway GR, Webster MA, Smith SM, Nichols TE (2014) Permutation inference for the general linear model. *NeuroImage* 92:381-397. doi:10.1016/j.neuroimage.2014.01.060